

R Exercise 3

Table of contents

1	Classification: Basic Concepts and Techniques	2
1.1	Install packages	2
1.2	The Zoo Dataset	2
1.3	Decision Trees	4
1.3.1	Create Tree With Default Settings (uses pre-pruning)	4
1.3.2	Create a Full Tree	6
1.3.3	Make Predictions for New Data	10
1.4	Model Evaluation with Caret	11
1.4.1	Hold out Test Data	11
1.4.2	Learn a Model and Tune Hyperparameters on the Training Data	11
1.5	Testing: Confusion Matrix and Confidence Interval for Accuracy . . .	14
1.6	Model Comparison	16
1.7	Feature Selection and Feature Preparation	18
1.7.1	Feature Subset Selection	21
1.7.2	Using Dummy Variables for Factors	22
1.8	Class Imbalance	24
1.8.1	Option 1: Use the Data As Is and Hope For The Best	27
1.8.2	Option 2: Balance Data With Resampling	29
1.8.3	Option 3: Build A Larger Tree and use Predicted Probabilities	32
1.8.4	Option 4: Use a Cost-Sensitive Classifier	36
2	Classification: Alternative Techniques	38
2.1	Install packages	38
2.2	Training and Test Data	39
2.3	Fitting Different Classification Models to the Training Data	40
2.3.1	Conditional Inference Tree (Decision Tree)	40
2.3.2	C 4.5 Decision Tree	41
2.3.3	K-Nearest Neighbors	43
2.3.4	PART (Rule-based classifier)	44

2.3.5	Linear Support Vector Machines	46
2.3.6	Random Forest	47
2.3.7	Gradient Boosted Decision Trees (xgboost)	48
2.3.8	Artificial Neural Network	50
2.4	Comparing Models	51
2.5	Applying the Chosen Model to the Test Data	55
2.6	Comparing Decision Boundaries of Popular Classification Techniques	56
2.6.1	Penguins Dataset	58
2.6.2	Circle Dataset	77
2.7	More Information on Classification with R	92

1 Classification: Basic Concepts and Techniques

1.1 Install packages

```
if(!require(pacman))
  install.packages("pacman")
```

Loading required package: pacman

```
pacman::p_load(tidyverse, rpart, rpart.plot, caret,
  lattice, FSelector, sampling, pROC, mlbench)
```

1.2 The Zoo Dataset

```
data(Zoo, package="mlbench")
head(Zoo)
```

	hair	feathers	eggs	milk	airborne	aquatic	predator	toothed	backbone
aardvark	TRUE	FALSE	FALSE	TRUE	FALSE	FALSE	TRUE	TRUE	TRUE
antelope	TRUE	FALSE	FALSE	TRUE	FALSE	FALSE	FALSE	TRUE	TRUE
bass	FALSE	FALSE	TRUE	FALSE	FALSE	TRUE	TRUE	TRUE	TRUE
bear	TRUE	FALSE	FALSE	TRUE	FALSE	FALSE	TRUE	TRUE	TRUE
boar	TRUE	FALSE	FALSE	TRUE	FALSE	FALSE	TRUE	TRUE	TRUE
buffalo	TRUE	FALSE	FALSE	TRUE	FALSE	FALSE	FALSE	TRUE	TRUE
	breathes	venomous	fins	legs	tail	domestic	catsize	type	
aardvark	TRUE	FALSE	FALSE	4	FALSE	FALSE	TRUE	mammal	

antelope	TRUE	FALSE	FALSE	4	TRUE	FALSE	TRUE	mammal
bass	FALSE	FALSE	TRUE	0	TRUE	FALSE	FALSE	fish
bear	TRUE	FALSE	FALSE	4	FALSE	FALSE	TRUE	mammal
boar	TRUE	FALSE	FALSE	4	TRUE	FALSE	TRUE	mammal
buffalo	TRUE	FALSE	FALSE	4	TRUE	FALSE	TRUE	mammal

```
library(tidyverse)
as_tibble(Zoo, rownames = "animal")
```

```
# A tibble: 101 x 18
  animal hair feathers eggs milk airborne aquatic predator toothed backbone
  <chr>   <lgl> <lgl>   <lgl> <lgl> <lgl>   <lgl>   <lgl>   <lgl>   <lgl>
1 aardva~ TRUE  FALSE   FALSE TRUE  FALSE   FALSE   TRUE    TRUE    TRUE
2 antelo~ TRUE  FALSE   FALSE TRUE  FALSE   FALSE   FALSE   TRUE    TRUE
3 bass    FALSE FALSE   TRUE  FALSE FALSE   TRUE    TRUE    TRUE    TRUE
4 bear    TRUE  FALSE   FALSE TRUE  FALSE   FALSE   TRUE    TRUE    TRUE
5 boar    TRUE  FALSE   FALSE TRUE  FALSE   FALSE   TRUE    TRUE    TRUE
6 buffalo TRUE  FALSE   FALSE TRUE  FALSE   FALSE   FALSE   TRUE    TRUE
7 calf    TRUE  FALSE   FALSE TRUE  FALSE   FALSE   FALSE   TRUE    TRUE
8 carp    FALSE FALSE   TRUE  FALSE FALSE   TRUE    FALSE   TRUE    TRUE
9 catfish FALSE FALSE   TRUE  FALSE FALSE   TRUE    TRUE    TRUE    TRUE
10 cavy    TRUE  FALSE   FALSE TRUE  FALSE   FALSE   FALSE   TRUE    TRUE
# i 91 more rows
# i 8 more variables: breathes <lgl>, venomous <lgl>, fins <lgl>, legs <int>,
#   tail <lgl>, domestic <lgl>, catsize <lgl>, type <fct>
```

```
Zoo <- Zoo |>
  mutate(across(where(is.logical), factor, levels = c(TRUE, FALSE))) |>
  mutate(across(where(is.character), factor))
```

Warning: There was 1 warning in `mutate()`.

i In argument: `across(where(is.logical), factor, levels = c(TRUE, FALSE))`.

Caused by warning:

! The `...` argument of `across()` is deprecated as of dplyr 1.1.0.

Supply arguments directly to `.fns` through an anonymous function instead.

```
# Previously
across(a:b, mean, na.rm = TRUE)

# Now
across(a:b, \(x) mean(x, na.rm = TRUE))
```

```
summary(Zoo)
```

hair	feathers	eggs	milk	airborne	aquatic	predator
TRUE :43	TRUE :20	TRUE :59	TRUE :41	TRUE :24	TRUE :36	TRUE :56
FALSE:58	FALSE:81	FALSE:42	FALSE:60	FALSE:77	FALSE:65	FALSE:45

toothed	backbone	breathes	venomous	fins	legs
TRUE :61	TRUE :83	TRUE :80	TRUE : 8	TRUE :17	Min. :0.000
FALSE:40	FALSE:18	FALSE:21	FALSE:93	FALSE:84	1st Qu.:2.000
					Median :4.000
					Mean :2.842
					3rd Qu.:4.000
					Max. :8.000

tail	domestic	catsize		type
TRUE :75	TRUE :13	TRUE :44	mammal	:41
FALSE:26	FALSE:88	FALSE:57	bird	:20
			reptile	: 5
			fish	:13
			amphibian	: 4
			insect	: 8
			mollusc.et.al	:10

1.3 Decision Trees

```
library(rpart)
```

1.3.1 Create Tree With Default Settings (uses pre-pruning)

```
tree_default <- Zoo |>
  rpart(type ~ ., data = _)
tree_default
```

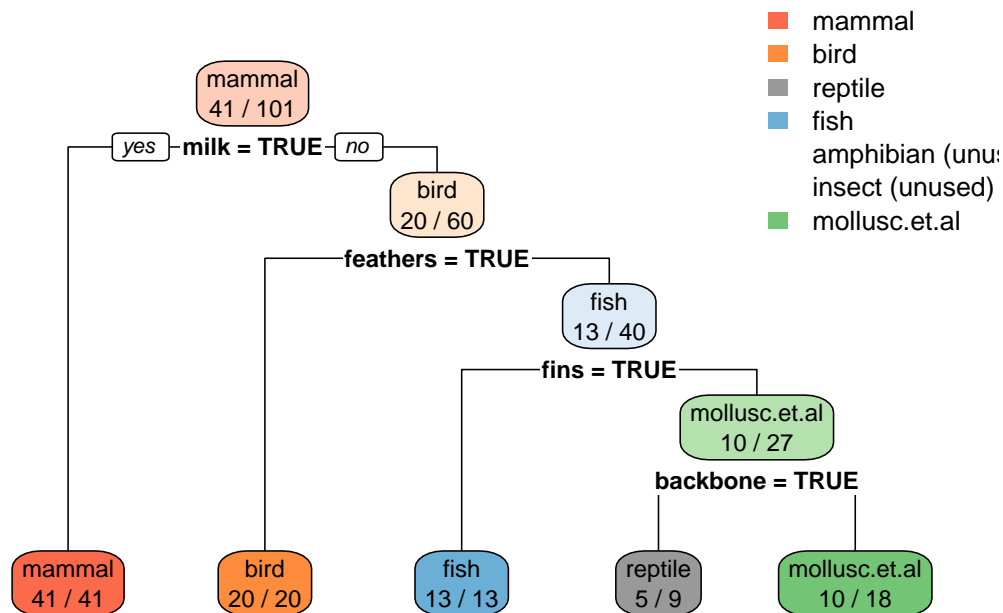
n= 101

```
node), split, n, loss, yval, (yprob)
  * denotes terminal node
```

```
1) root 101 60 mammal (0.41 0.2 0.05 0.13 0.04 0.079 0.099)
  2) milk=TRUE 41 0 mammal (1 0 0 0 0 0 0) *
  3) milk=FALSE 60 40 bird (0 0.33 0.083 0.22 0.067 0.13 0.17)
    6) feathers=TRUE 20 0 bird (0 1 0 0 0 0 0) *
    7) feathers=FALSE 40 27 fish (0 0 0.12 0.33 0.1 0.2 0.25)
      14) fins=TRUE 13 0 fish (0 0 0 1 0 0 0) *
      15) fins=FALSE 27 17 mollusc.et.al (0 0 0.19 0 0.15 0.3 0.37)
        30) backbone=TRUE 9 4 reptile (0 0 0.56 0 0.44 0 0) *
        31) backbone=FALSE 18 8 mollusc.et.al (0 0 0 0 0 0.44 0.56) *
```

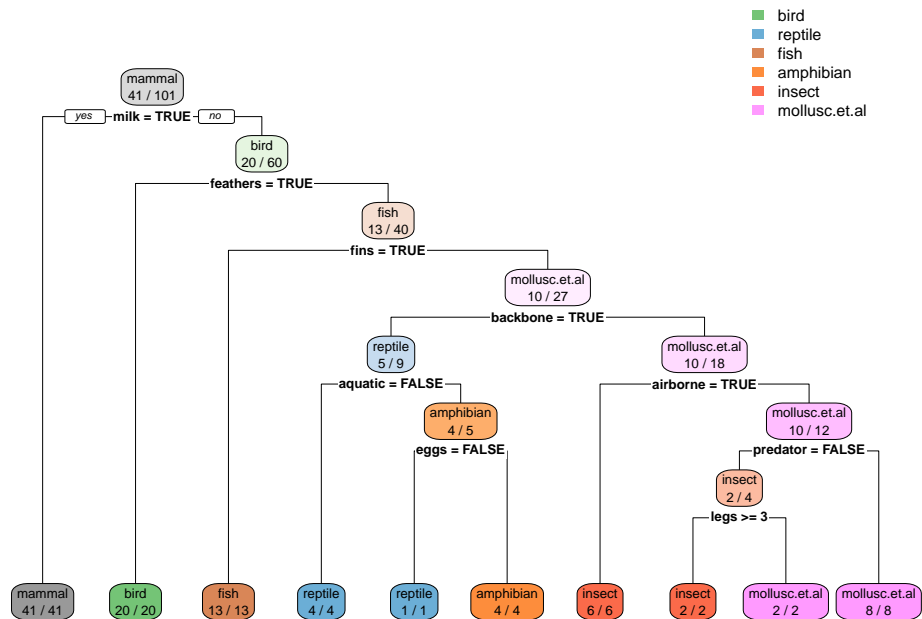
Plotting

```
library(rpart.plot)
rpart.plot(tree_default, extra = 2)
```



1.3.2 Create a Full Tree

```
tree_full <- Zoo |>
  rpart(type ~ . , data = _,
        control = rpart.control(minsplit = 2, cp = 0))
rpart.plot(tree_full, extra = 2,
           roundint=FALSE,
           box.palette = list("Gy", "Gn", "Bu", "Bn",
                             "Or", "Rd", "Pu")) # specify 7 colors
```



tree_full

n= 101

node), split, n, loss, yval, (yprob)
* denotes terminal node

```
1) root 101 60 mammal (0.41 0.2 0.05 0.13 0.04 0.079 0.099)
  2) milk=TRUE 41 0 mammal (1 0 0 0 0 0 0) *
  3) milk=FALSE 60 40 bird (0 0.33 0.083 0.22 0.067 0.13 0.17)
    6) feathers=TRUE 20 0 bird (0 1 0 0 0 0 0) *
    7) feathers=FALSE 40 27 fish (0 0 0.12 0.33 0.1 0.2 0.25)
```

```

14) fins=TRUE 13 0 fish (0 0 0 1 0 0 0) *
15) fins=FALSE 27 17 mollusc.et.al (0 0 0.19 0 0.15 0.3 0.37)
30) backbone=TRUE 9 4 reptile (0 0 0.56 0 0.44 0 0)
60) aquatic=FALSE 4 0 reptile (0 0 1 0 0 0 0) *
61) aquatic=TRUE 5 1 amphibian (0 0 0.2 0 0.8 0 0)
122) eggs=FALSE 1 0 reptile (0 0 1 0 0 0 0) *
123) eggs=TRUE 4 0 amphibian (0 0 0 0 1 0 0) *
31) backbone=FALSE 18 8 mollusc.et.al (0 0 0 0 0 0.44 0.56)
62) airborne=TRUE 6 0 insect (0 0 0 0 0 1 0) *
63) airborne=FALSE 12 2 mollusc.et.al (0 0 0 0 0 0.17 0.83)
126) predator=FALSE 4 2 insect (0 0 0 0 0 0.5 0.5)
252) legs>=3 2 0 insect (0 0 0 0 0 1 0) *
253) legs< 3 2 0 mollusc.et.al (0 0 0 0 0 0 1) *
127) predator=TRUE 8 0 mollusc.et.al (0 0 0 0 0 0 1) *

```

Training error on tree with pre-pruning

```
predict(tree_default, Zoo) |> head ()
```

	mammal	bird	reptile	fish	amphibian	insect	mollusc.et.al
aardvark	1	0	0	0	0	0	0
antelope	1	0	0	0	0	0	0
bass	0	0	0	1	0	0	0
bear	1	0	0	0	0	0	0
boar	1	0	0	0	0	0	0
buffalo	1	0	0	0	0	0	0

```

pred <- predict(tree_default, Zoo, type="class")
head(pred)

```

```

aardvark antelope    bass    bear    boar  buffalo
mammal    mammal    fish    mammal  mammal  mammal
Levels: mammal bird reptile fish amphibian insect mollusc.et.al

```

```

confusion_table <- with(Zoo, table(type, pred))
confusion_table

```

	pred
type	mammal bird reptile fish amphibian insect mollusc.et.al

mammal	41	0	0	0	0	0	0
bird	0	20	0	0	0	0	0
reptile	0	0	5	0	0	0	0
fish	0	0	0	13	0	0	0
amphibian	0	0	4	0	0	0	0
insect	0	0	0	0	0	0	8
mollusc.et.al	0	0	0	0	0	0	10

```
correct <- confusion_table |> diag() |> sum()
correct
```

[1] 89

```
error <- confusion_table |> sum() - correct
error
```

[1] 12

```
accuracy <- correct / (correct + error)
accuracy
```

[1] 0.8811881

Use a function for accuracy

```
accuracy <- function(truth, prediction) {
  tbl <- table(truth, prediction)
  sum(diag(tbl))/sum(tbl)
}
```

```
accuracy(Zoo |> pull(type), pred)
```

[1] 0.8811881

Training error of the full tree


```
accuracy(Zoo |> pull(type),
         predict(tree_full, Zoo, type = "class"))
```

```
[1] 1
```

Get a confusion table with more statistics (using caret)

```
library(caret)
confusionMatrix(data = pred,
                 reference = Zoo |> pull(type))
```

Confusion Matrix and Statistics

	Reference						
Prediction	mammal	bird	reptile	fish	amphibian	insect	mollusc.et.al
mammal	41	0	0	0	0	0	0
bird	0	20	0	0	0	0	0
reptile	0	0	5	0	4	0	0
fish	0	0	0	13	0	0	0
amphibian	0	0	0	0	0	0	0
insect	0	0	0	0	0	0	0
mollusc.et.al	0	0	0	0	0	8	10

Overall Statistics

```

Accuracy : 0.8812
 95% CI : (0.8017, 0.9371)
No Information Rate : 0.4059
P-Value [Acc > NIR] : < 2.2e-16
```

```
Kappa : 0.8431
```

```
McNemar's Test P-Value : NA
```

Statistics by Class:

	Class: mammal	Class: bird	Class: reptile	Class: fish
Sensitivity	1.0000	1.000	1.00000	1.0000
Specificity	1.0000	1.000	0.95833	1.0000
Pos Pred Value	1.0000	1.000	0.55556	1.0000

Neg Pred Value	1.0000	1.000	1.00000	1.0000
Prevalence	0.4059	0.198	0.04950	0.1287
Detection Rate	0.4059	0.198	0.04950	0.1287
Detection Prevalence	0.4059	0.198	0.08911	0.1287
Balanced Accuracy	1.0000	1.000	0.97917	1.0000
	Class: amphibian	Class: insect	Class: mollusc.et.al	
Sensitivity	0.0000	0.00000		1.00000
Specificity	1.0000	1.00000		0.91209
Pos Pred Value	NaN	NaN		0.55556
Neg Pred Value	0.9604	0.92079		1.00000
Prevalence	0.0396	0.07921		0.09901
Detection Rate	0.0000	0.00000		0.09901
Detection Prevalence	0.0000	0.00000		0.17822
Balanced Accuracy	0.5000	0.50000		0.95604

1.3.3 Make Predictions for New Data

```
my_animal <- tibble(hair = TRUE, feathers = TRUE, eggs = FALSE,
  milk = TRUE, airborne = TRUE, aquatic = FALSE, predator = TRUE,
  toothed = TRUE, backbone = TRUE, breathes = TRUE, venomous = FALSE,
  fins = FALSE, legs = 4, tail = TRUE, domestic = FALSE,
  catsize = FALSE, type = NA)
```

```
my_animal <- my_animal |>
  mutate(across(where(is.logical), factor, levels = c(TRUE, FALSE)))
my_animal
```

```
# A tibble: 1 x 17
  hair feathers eggs milk airborne aquatic predator toothed backbone breathes
<fct> <fct>   <fct> <fct> <fct>   <fct>   <fct>   <fct>   <fct>   <fct>
1 TRUE  TRUE    FALSE TRUE  TRUE    FALSE   TRUE    TRUE    TRUE    TRUE
# i 7 more variables: venomous <fct>, fins <fct>, legs <dbl>, tail <fct>,
# domestic <fct>, catsize <fct>, type <fct>
```

```
predict(tree_default , my_animal, type = "class")
```

```
1
mammal
Levels: mammal bird reptile fish amphibian insect mollusc.et.al
```

1.4 Model Evaluation with Caret

```
library(caret)
```

Set random number generator seed to make results reproducible

```
set.seed(2000)
```

1.4.1 Hold out Test Data

```
inTrain <- createDataPartition(y = Zoo$type, p = .8, list = FALSE)
Zoo_train <- Zoo |> slice(inTrain)
```

Warning: Slicing with a 1-column matrix was deprecated in dplyr 1.1.0.

```
Zoo_test <- Zoo |> slice(-inTrain)
```

1.4.2 Learn a Model and Tune Hyperparameters on the Training Data

```
fit <- Zoo_train |>
  train(type ~ .,
        data = _ ,
        method = "rpart",
        control = rpart.control(minsplit = 2),
        trControl = trainControl(method = "cv", number = 10),
        tuneLength = 5)

fit
```

CART

83 samples

16 predictors

7 classes: 'mammal', 'bird', 'reptile', 'fish', 'amphibian', 'insect', 'mollusc.etal'

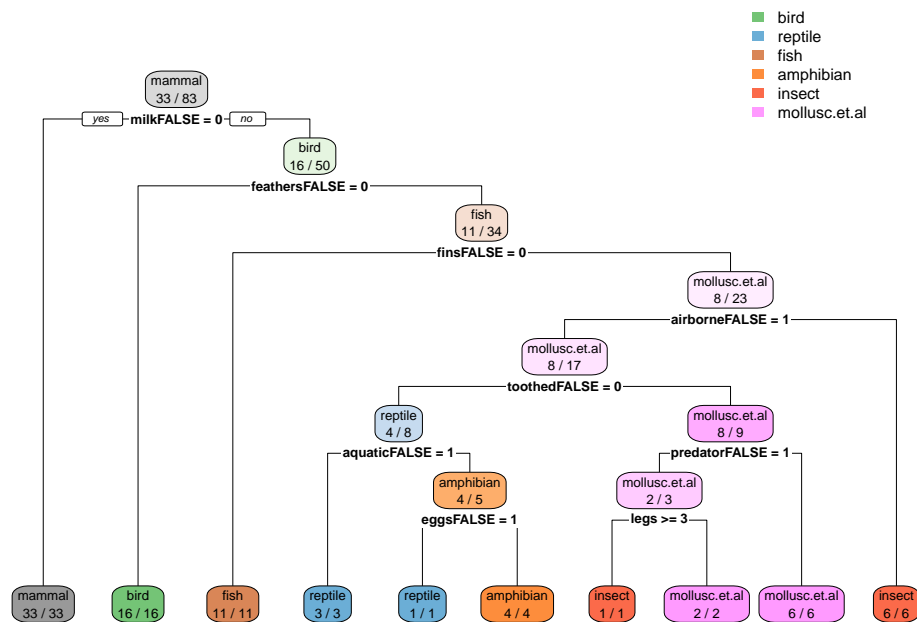
No pre-processing

Resampling: Cross-Validated (10 fold)

Resampling results across tuning parameters:

cp	Accuracy	Kappa
0.00	0.9384921	0.9188571
0.08	0.8973810	0.8681837
0.16	0.7447619	0.6637060
0.22	0.6663095	0.5540490
0.32	0.4735317	0.1900043

```
rpart.plot(fit$finalModel, extra = 2,
  box.palette = list("Gy", "Gn", "Bu", "Bn", "Or", "Rd", "Pu"))
```



```
varImp(fit)
```

```
rpart variable importance
```

	Overall
toothedFALSE	100.000


feathersFALSE	69.814
backboneFALSE	63.084
milkFALSE	55.555
eggsFALSE	53.614
hairFALSE	50.518
finsFALSE	46.984
tailFALSE	28.447
breathesFALSE	28.128
airborneFALSE	26.272
legs	25.859
aquaticFALSE	5.960
predatorFALSE	2.349
venomousFALSE	1.387
catsizeFALSE	0.000
domesticFALSE	0.000

```
imp <- varImp(fit, compete = FALSE)
imp
```

rpart variable importance

	Overall
milkFALSE	100.000
feathersFALSE	55.694
finsFALSE	39.453
toothedFALSE	22.956
airborneFALSE	22.478
aquaticFALSE	9.987
eggsFALSE	6.658
legs	5.549
predatorFALSE	1.850
domesticFALSE	0.000
breathesFALSE	0.000
catsizeFALSE	0.000
tailFALSE	0.000
hairFALSE	0.000
backboneFALSE	0.000
venomousFALSE	0.000

```
ggplot(imp)
```



1.5 Testing: Confusion Matrix and Confidence Interval for Accuracy

Use the best model on the test data

```
pred <- predict(fit, newdata = Zoo_test)
pred
```

```
[1] mammal      mammal      mollusc.et.al insect      mammal
[6] mammal      mammal      bird        mammal      mammal
[11] bird        fish        fish        mammal      mollusc.et.al
[16] bird        insect      bird

Levels: mammal bird reptile fish amphibian insect mollusc.et.al
```

```
confusionMatrix(data = pred,
                  ref = Zoo_test |> pull(type))
```

Confusion Matrix and Statistics

	Reference						
Prediction	mammal	bird	reptile	fish	amphibian	insect	mollusc.et.al
mammal	8	0	0	0	0	0	0

bird	0	4	0	0	0	0	0
reptile	0	0	0	0	0	0	0
fish	0	0	0	2	0	0	0
amphibian	0	0	0	0	0	0	0
insect	0	0	1	0	0	1	0
mollusc.et.al	0	0	0	0	0	0	2

Overall Statistics

Accuracy : 0.9444
 95% CI : (0.7271, 0.9986)
 No Information Rate : 0.4444
 P-Value [Acc > NIR] : 1.076e-05

Kappa : 0.9231

McNemar's Test P-Value : NA

Statistics by Class:

	Class: mammal	Class: bird	Class: reptile	Class: fish
Sensitivity	1.0000	1.0000	0.00000	1.0000
Specificity	1.0000	1.0000	1.00000	1.0000
Pos Pred Value	1.0000	1.0000	NaN	1.0000
Neg Pred Value	1.0000	1.0000	0.94444	1.0000
Prevalence	0.4444	0.2222	0.05556	0.1111
Detection Rate	0.4444	0.2222	0.00000	0.1111
Detection Prevalence	0.4444	0.2222	0.00000	0.1111
Balanced Accuracy	1.0000	1.0000	0.50000	1.0000
	Class: amphibian	Class: insect	Class: mollusc.et.al	
Sensitivity	NA	1.00000	1.0000	
Specificity	1	0.94118	1.0000	
Pos Pred Value	NA	0.50000	1.0000	
Neg Pred Value	NA	1.00000	1.0000	
Prevalence	0	0.05556	0.1111	
Detection Rate	0	0.05556	0.1111	
Detection Prevalence	0	0.11111	0.1111	
Balanced Accuracy	NA	0.97059	1.0000	

1.6 Model Comparison

```
train_index <- createFolds(Zoo_train$type, k = 10)

rpartFit <- Zoo_train |>
  train(type ~ .,
        data = _,
        method = "rpart",
        tuneLength = 10,
        trControl = trainControl(method = "cv", indexOut = train_index)
  )

knnFit <- Zoo_train |>
  train(type ~ .,
        data = _,
        method = "knn",
        preProcess = "scale",
        tuneLength = 10,
        trControl = trainControl(method = "cv", indexOut = train_index)
  )

resamps <- resamples(list(
  CART = rpartFit,
  kNearestNeighbors = knnFit
))

summary(resamps)
```

Call:

```
summary.resamples(object = resamps)
```

Models: CART, kNearestNeighbors

Number of resamples: 10

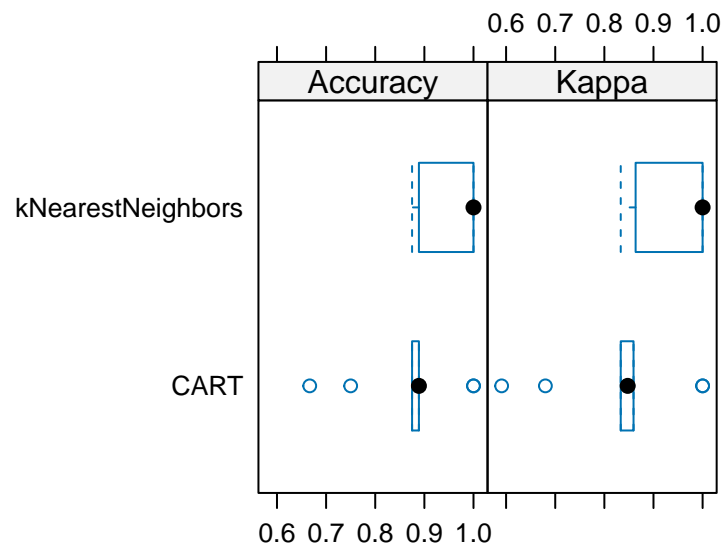
Accuracy

	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.	NA's
CART	0.6666667	0.8750000	0.8888889	0.8722222	0.8888889	1	0
kNearestNeighbors	0.8750000	0.9166667	1.0000000	0.9652778	1.0000000	1	0

Kappa

	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.	NA's
CART	0.5909091	0.8333333	0.8474576	0.8341866	0.8570312	1	0
kNearestNeighbors	0.8333333	0.8977273	1.0000000	0.9546970	1.0000000	1	0

```
library(lattice)
bwplot(resamps, layout = c(3, 1))
```



```
difs <- diff(resamps)
difs
```

Call:
diff.resamples(x = resamps)

Models: CART, kNearestNeighbors
Metrics: Accuracy, Kappa
Number of differences: 1
p-value adjustment: bonferroni

```
summary(difs)
```

Call:
summary.diff.resamples(object = difs)

p-value adjustment: bonferroni
Upper diagonal: estimates of the difference
Lower diagonal: p-value for H0: difference = 0

Accuracy

	CART	kNearestNeighbors
CART		-0.09306
kNearestNeighbors	0.01151	

Kappa

	CART	kNearestNeighbors
CART		-0.1205
kNearestNeighbors	0.0104	

1.7 Feature Selection and Feature Preparation

```
library(FSelector)
```


```
weights <- Zoo_train |>  
  chi.squared(type ~ ., data = _) |>  
  as_tibble(rownames = "feature") |>  
  arrange(desc(attr_importance))
```

```
weights
```

```
# A tibble: 16 x 2  
  feature attr_importance  
  <chr>      <dbl>  
1 feathers      1  
2 backbone      1  
3 milk          1  
4 toothed    0.975  
5 eggs        0.933  
6 hair        0.907  
7 breathes    0.898  
8 airborne    0.848  
9 fins        0.845  
10 legs       0.828  
11 tail       0.779  
12 catsize    0.664
```

13	aquatic	0.655
14	venomous	0.475
15	predator	0.385
16	domestic	0.231

```
ggplot(weights,
  aes(x = attr_importance, y = reorder(feature, attr_importance))) +
  geom_bar(stat = "identity") +
  xlab("Importance score") +
  ylab("Feature")
```



```
# Get the 5 best features
```

```
subset <- cutoff.k(weights |>
  column_to_rownames("feature"), 5)
subset
```

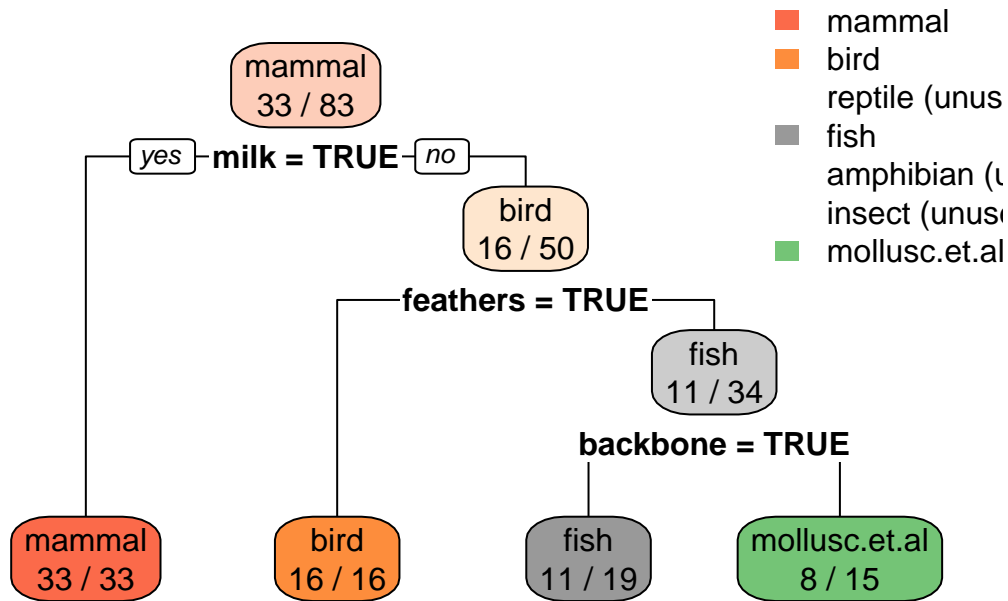
```
[1] "feathers" "backbone" "milk"      "toothed"  "eggs"
```

```
# Use only the best 5 features to build a model (Fselector provides as.simple.formula)
```

```
f <- as.simple.formula(subset, "type")
f
```

```
type ~ feathers + backbone + milk + toothed + eggs
<environment: 0x121de8da0>
```

```
m <- Zoo_train |> rpart(f, data = _)
rpart.plot(m, extra = 2, roundint = FALSE)
```



```
Zoo_train |>
  gain.ratio(type ~ ., data = _) |>
  as_tibble(rownames = "feature") |>
  arrange(desc(attr_importance))
```

```
# A tibble: 16 x 2
  feature attr_importance
  <chr>      <dbl>
1 backbone      1
2 milk          1.00
3 feathers      1.00
```

4	toothed	0.919
5	eggs	0.827
6	breathes	0.821
7	hair	0.782
8	fins	0.689
9	legs	0.682
10	airborne	0.671
11	tail	0.573
12	aquatic	0.391
13	catsize	0.383
14	venomous	0.351
15	predator	0.125
16	domestic	0.0975

1.7.1 Feature Subset Selection

```
Zoo_train |>
  cfs(type ~ ., data = _)
```

```
[1] "hair"      "feathers" "eggs"      "milk"      "toothed"   "backbone"
[7] "breathes" "fins"      "legs"      "tail"
```


```
evaluator <- function(subset) {
  model <- Zoo_train |>
    train(as.simple.formula(subset, "type"),
          data = _,
          method = "rpart",
          trControl = trainControl(method = "boot", number = 5),
          tuneLength = 0)
  results <- model$resample$Accuracy
  cat("Trying features:", paste(subset, collapse = " + "), "\n")
  m <- mean(results)
  cat("Accuracy:", round(m, 2), "\n\n")
  m
}
```

```
features <- Zoo_train |> colnames() |> setdiff("type")
```

```
##subset <- backward.search(features, evaluator)
##subset <- forward.search(features, evaluator)
##subset <- best.first.search(features, evaluator)
##subset <- hill.climbing.search(features, evaluator)
##subset
```

1.7.2 Using Dummy Variables for Factors

```
tree_predator <- Zoo_train |>
  rpart(predator ~ type, data = _)
rpart.plot(tree_predator, extra = 2, roundint = FALSE)
```



```
Zoo_train_dummy <- as_tibble(class2ind(Zoo_train$type)) |>
  mutate(across(everything(), as.factor)) |>
  add_column(predator = Zoo_train$predator)
Zoo_train_dummy
```

A tibble: 83 x 8

	mammal	bird	reptile	fish	amphibian	insect	mollusc.et.al	predator
	<fct>	<fct>	<fct>	<fct>	<fct>	<fct>	<fct>	<fct>
1	1	0	0	0	0	0	0	TRUE
2	1	0	0	0	0	0	0	FALSE
3	0	0	0	1	0	0	0	TRUE
4	1	0	0	0	0	0	0	TRUE
5	1	0	0	0	0	0	0	FALSE
6	1	0	0	0	0	0	0	FALSE
7	0	0	0	1	0	0	0	FALSE
8	0	0	0	1	0	0	0	TRUE

```

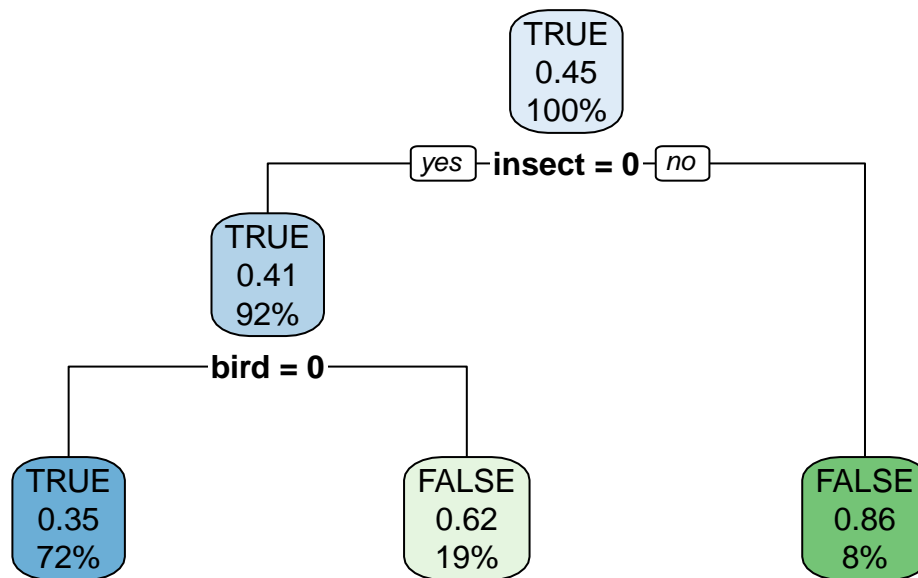
 9 1      0      0      0      0      0      0      FALSE
10 0      1      0      0      0      0      0      FALSE
# i 73 more rows

```

```

tree_predator <- Zoo_train_dummy |>
  rpart(predator ~ .,
    data = _,
    control = rpart.control(minsplit = 2, cp = 0.01))
rpart.plot(tree_predator, roundint = FALSE)

```



```

fit <- Zoo_train |>
  train(predator ~ type,
    data = _,
    method = "rpart",
    control = rpart.control(minsplit = 2),
    tuneGrid = data.frame(cp = 0.01))
fit

```

CART

```

83 samples
1 predictor
2 classes: 'TRUE', 'FALSE'

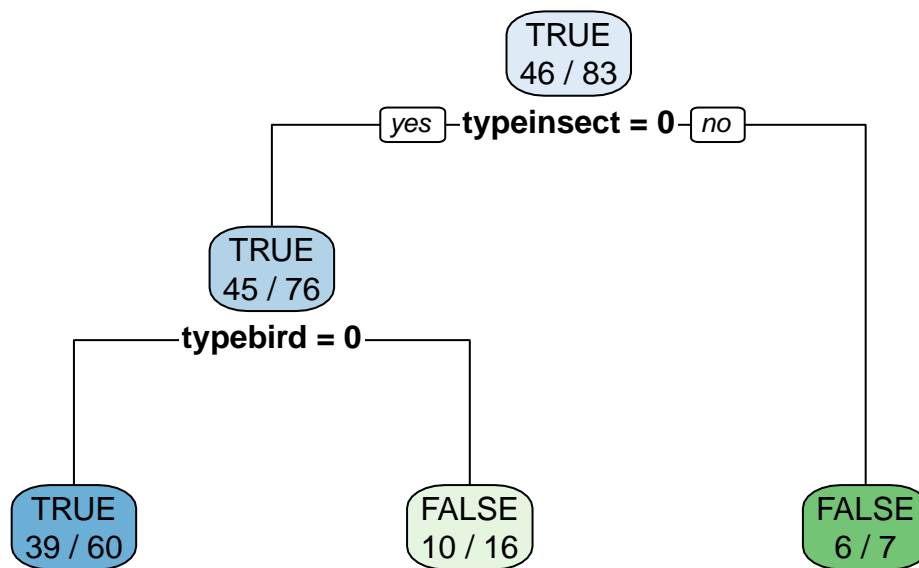
```

No pre-processing
Resampling: Bootstrapped (25 reps)
Summary of sample sizes: 83, 83, 83, 83, 83, 83, ...
Resampling results:

Accuracy	Kappa
0.6060423	0.2034198

Tuning parameter 'cp' was held constant at a value of 0.01

```
rpart.plot(fit$finalModel, extra = 2)
```




1.8 Class Imbalance

```
library(rpart)
library(rpart.plot)
data(Zoo, package="mlbench")

#Class distribution

ggplot(Zoo, aes(y = type)) + geom_bar()
```

```
Zoo_reptile <- Zoo |>
  mutate(type = factor(Zoo$type == "reptile",
    levels = c(FALSE, TRUE),
    labels = c("nonreptile", "reptile")))
```

```
summary(Zoo_reptile)
```

hair	feathers	eggs	milk
Mode :logical	Mode :logical	Mode :logical	Mode :logical
FALSE:58	FALSE:81	FALSE:42	FALSE:60
TRUE :43	TRUE :20	TRUE :59	TRUE :41


airborne	aquatic	predator	toothed
Mode :logical	Mode :logical	Mode :logical	Mode :logical
FALSE:77	FALSE:65	FALSE:45	FALSE:40
TRUE :24	TRUE :36	TRUE :56	TRUE :61

backbone	breathes	venomous	fins
----------	----------	----------	------

Mode :logical	Mode :logical	Mode :logical	Mode :logical
FALSE:18	FALSE:21	FALSE:93	FALSE:84
TRUE :83	TRUE :80	TRUE :8	TRUE :17

legs	tail	domestic	catsize
Min. :0.000	Mode :logical	Mode :logical	Mode :logical
1st Qu.:2.000	FALSE:26	FALSE:88	FALSE:57
Median :4.000	TRUE :75	TRUE :13	TRUE :44
Mean :2.842			
3rd Qu.:4.000			
Max. :8.000			
type			
nonreptile:96			
reptile : 5			

```
ggplot(Zoo_reptile, aes(y = type)) + geom_bar()
```



```

set.seed(1234)

inTrain <- createDataPartition(y = Zoo_reptile$type, p = .5, list = FALSE)
training_reptile <- Zoo_reptile |> slice(inTrain)
testing_reptile <- Zoo_reptile |> slice(-inTrain)

```

1.8.1 Option 1: Use the Data As Is and Hope For The Best

```

fit <- training_reptile |>
  train(type ~ .,
        data = _,
        method = "rpart",
        trControl = trainControl(method = "cv"))

```

Warning in nominalTrainWorkflow(x = x, y = y, wts = weights, info = trainInfo,
: There were missing values in resampled performance measures.

```
fit
```

CART

51 samples
16 predictors
2 classes: 'nonreptile', 'reptile'

No pre-processing
Resampling: Cross-Validated (10 fold)
Summary of sample sizes: 46, 47, 46, 46, 45, 46, ...
Resampling results:

Accuracy	Kappa
0.9466667	0

Tuning parameter 'cp' was held constant at a value of 0

```
rpart.plot(fit$finalModel, extra = 2)
```

```
confusionMatrix(data = predict(fit, testing_reptile),
  ref = testing_reptile$type, positive = "reptile")
```

Confusion Matrix and Statistics

	Reference	
Prediction	nonreptile	reptile
nonreptile	48	2
reptile	0	0

```

      Accuracy : 0.96
      95% CI   : (0.8629, 0.9951)
No Information Rate : 0.96
P-Value [Acc > NIR] : 0.6767

```

```
Kappa : 0
```

```
McNemar's Test P-Value : 0.4795
```

```

      Sensitivity : 0.00
      Specificity : 1.00
Pos Pred Value   : NaN
Neg Pred Value   : 0.96
      Prevalence  : 0.04
      Detection Rate : 0.00
Detection Prevalence : 0.00
Balanced Accuracy : 0.50

```

```
'Positive' Class : reptile
```

1.8.2 Option 2: Balance Data With Resampling

```
library(sampling)
set.seed(1000) # for repeatability

id <- strata(training_reptile, stratanames = "type", size = c(50, 50), method = "srswr")
training_reptile_balanced <- training_reptile |>
  slice(id$ID_unit)
table(training_reptile_balanced$type)
```

```
nonreptile    reptile
          50         50
```

```
fit <- training_reptile_balanced |>
  train(type ~ .,
        data = _,
        method = "rpart",
        trControl = trainControl(method = "cv"),
        control = rpart.control(minsplit = 5))

fit
```

CART

```
100 samples
16 predictor
2 classes: 'nonreptile', 'reptile'
```

No pre-processing

Resampling: Cross-Validated (10 fold)

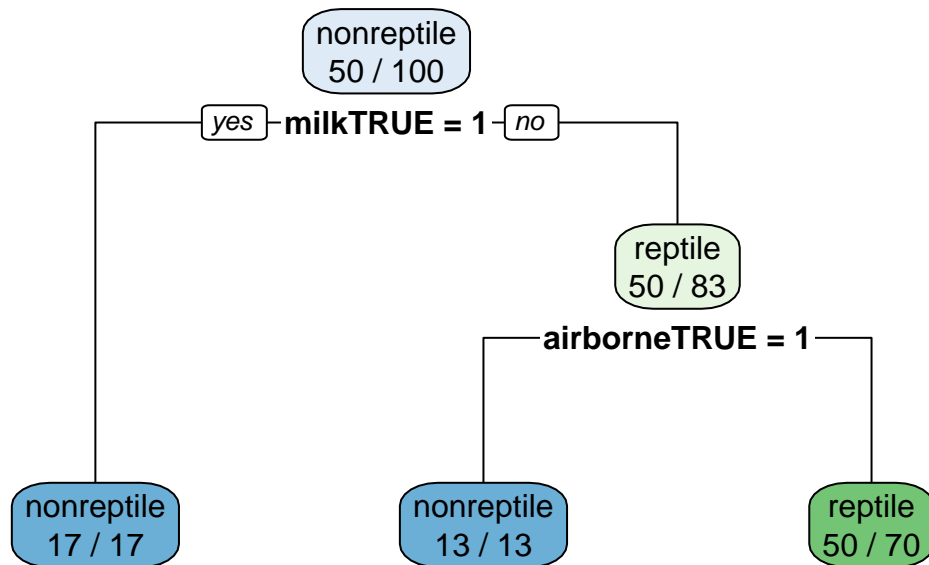
Summary of sample sizes: 90, 90, 90, 90, 90, 90, ...

Resampling results across tuning parameters:

cp	Accuracy	Kappa
0.22	0.78	0.56
0.26	0.67	0.34
0.34	0.54	0.08

Accuracy was used to select the optimal model using the largest value.
The final value used for the model was cp = 0.22.

```
rpart.plot(fit$finalModel, extra = 2)
```



```
confusionMatrix(data = predict(fit, testing_reptile),
  ref = testing_reptile$type, positive = "reptile")
```

Confusion Matrix and Statistics

	Reference	
Prediction	nonreptile	reptile
nonreptile	33	0
reptile	15	2

Accuracy : 0.7
 95% CI : (0.5539, 0.8214)
 No Information Rate : 0.96
 P-Value [Acc > NIR] : 1.0000000

Kappa : 0.1497

Mcnemar's Test P-Value : 0.0003006

Sensitivity : 1.0000
 Specificity : 0.6875
 Pos Pred Value : 0.1176

```

Neg Pred Value : 1.0000
Prevalence : 0.0400
Detection Rate : 0.0400
Detection Prevalence : 0.3400
Balanced Accuracy : 0.8438

'Positive' Class : reptile

```

```

id <- strata(training_reptile, stratanames = "type", size = c(50, 100), method = "srswr")
training_reptile_balanced <- training_reptile |>
  slice(id$ID_unit)
table(training_reptile_balanced$type)

```

```

nonreptile    reptile
      50         100

```

```

fit <- training_reptile_balanced |>
  train(type ~ .,
        data = _,
        method = "rpart",
        trControl = trainControl(method = "cv"),
        control = rpart.control(minsplit = 5))

confusionMatrix(data = predict(fit, testing_reptile),
                 ref = testing_reptile$type, positive = "reptile")

```

Confusion Matrix and Statistics

	Reference	
Prediction	nonreptile	reptile
nonreptile	33	0
reptile	15	2

```

Accuracy : 0.7
95% CI : (0.5539, 0.8214)
No Information Rate : 0.96
P-Value [Acc > NIR] : 1.0000000

```

```

Kappa : 0.1497

```

McNemar's Test P-Value : 0.0003006

Sensitivity : 1.0000
Specificity : 0.6875
Pos Pred Value : 0.1176
Neg Pred Value : 1.0000
Prevalence : 0.0400
Detection Rate : 0.0400
Detection Prevalence : 0.3400
Balanced Accuracy : 0.8438

'Positive' Class : reptile

1.8.3 Option 3: Build A Larger Tree and use Predicted Probabilities

```
fit <- training_reptile |>
  train(type ~ .,
        data = _,
        method = "rpart",
        tuneLength = 10,
        trControl = trainControl(method = "cv",
                                   classProbs = TRUE, ## necessary for predict with type="prob"
                                   summaryFunction=twoClassSummary), ## necessary for ROC
        metric = "ROC",
        control = rpart.control(minsplit = 3))
```

Warning in nominalTrainWorkflow(x = x, y = y, wts = weights, info = trainInfo,
: There were missing values in resampled performance measures.

```
fit
```

CART

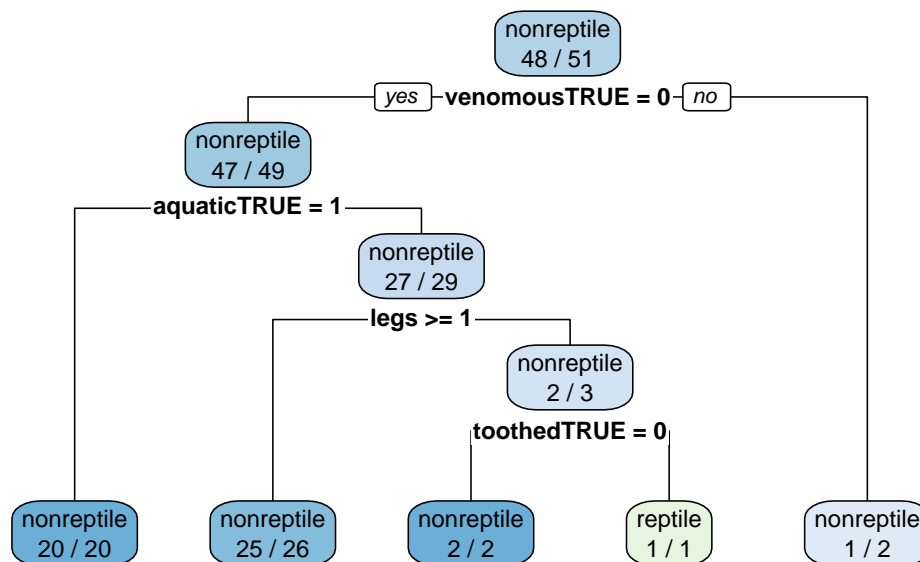
51 samples
16 predictors
2 classes: 'nonreptile', 'reptile'

No pre-processing
 Resampling: Cross-Validated (10 fold)
 Summary of sample sizes: 46, 47, 46, 46, 46, 45, ...
 Resampling results:

ROC	Sens	Spec
0.3583333	0.975	0

Tuning parameter 'cp' was held constant at a value of 0

```
rpart.plot(fit$finalModel, extra = 2)
```



```
confusionMatrix(data = predict(fit, testing_reptile),
  ref = testing_reptile$type, positive = "reptile")
```

Confusion Matrix and Statistics

	Reference	
Prediction	nonreptile	reptile
nonreptile	48	2
reptile	0	0

Accuracy : 0.96
 95% CI : (0.8629, 0.9951)

No Information Rate : 0.96
P-Value [Acc > NIR] : 0.6767

Kappa : 0

McNemar's Test P-Value : 0.4795

Sensitivity : 0.00
Specificity : 1.00
Pos Pred Value : NaN
Neg Pred Value : 0.96
Prevalence : 0.04
Detection Rate : 0.00
Detection Prevalence : 0.00
Balanced Accuracy : 0.50

'Positive' Class : reptile

```
prob <- predict(fit, testing_reptile, type = "prob")  
tail(prob)
```

	nonreptile	reptile
tuna	1.0000000	0.0000000
vole	0.9615385	0.03846154
wasp	0.5000000	0.5000000
wolf	0.9615385	0.03846154
worm	1.0000000	0.0000000
wren	0.9615385	0.03846154

```
pred <- as.factor(ifelse(prob[, "reptile"] >= 0.01, "reptile", "nonreptile"))  
  
confusionMatrix(data = pred,  
                 ref = testing_reptile$type, positive = "reptile")
```

Confusion Matrix and Statistics

	Reference	
Prediction	nonreptile	reptile
nonreptile	13	0

```

reptile          35          2

      Accuracy : 0.3
      95% CI   : (0.1786, 0.4461)
No Information Rate : 0.96
P-Value [Acc > NIR] : 1

```

```

      Kappa : 0.0289

```

```

McNemar's Test P-Value : 9.081e-09

```

```

      Sensitivity : 1.00000
      Specificity : 0.27083
      Pos Pred Value : 0.05405
      Neg Pred Value : 1.00000
      Prevalence : 0.04000
      Detection Rate : 0.04000
      Detection Prevalence : 0.74000
      Balanced Accuracy : 0.63542

```

```

'Positive' Class : reptile

```

```

library("pROC")
r <- roc(testing_reptile$type == "reptile", prob[, "reptile"])

```

```

Setting levels: control = FALSE, case = TRUE

```

```

Setting direction: controls < cases

```

```

r

```

```

Call:

```

```

roc.default(response = testing_reptile$type == "reptile", predictor = prob[, "reptile"])

```

```


Data: prob[, "reptile"] in 48 controls (testing_reptile$type == "reptile" FALSE) < 2 cases (1)
Area under the curve: 0.7656

```

```

ggroc(r) + geom_abline(intercept = 1, slope = 1, color = "darkgrey")

```



1.8.4 Option 4: Use a Cost-Sensitive Classifier

```
cost <- matrix(c(
  0, 1,
  100, 0
), byrow = TRUE, nrow = 2)
cost
```

```
      [,1] [,2]
[1,]    0    1
[2,]  100    0
```

```
fit <- training_reptile |>
  train(type ~ .,
        data = _,
        method = "rpart",
        parms = list(loss = cost),
        trControl = trainControl(method = "cv"))
```

```
fit
```

CART

51 samples

16 predictors

2 classes: 'nonreptile', 'reptile'

No pre-processing

Resampling: Cross-Validated (10 fold)

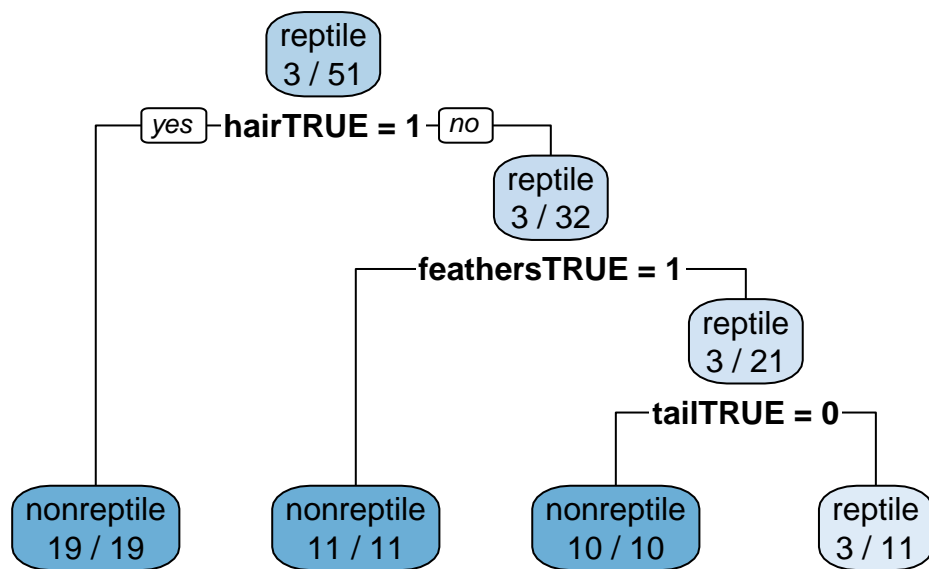
Summary of sample sizes: 46, 46, 46, 45, 46, 45, ...

Resampling results:

Accuracy	Kappa
0.4766667	-0.03038961

Tuning parameter 'cp' was held constant at a value of 0

```
rpart.plot(fit$finalModel, extra = 2)
```



```
confusionMatrix(data = predict(fit, testing_reptile),
  ref = testing_reptile$type, positive = "reptile")
```

Confusion Matrix and Statistics

	Reference	
Prediction	nonreptile	reptile
nonreptile	39	0
reptile	9	2

Accuracy : 0.82
 95% CI : (0.6856, 0.9142)
 No Information Rate : 0.96
 P-Value [Acc > NIR] : 0.999975

 Kappa : 0.2574

 McNemar's Test P-Value : 0.007661

 Sensitivity : 1.0000
 Specificity : 0.8125
 Pos Pred Value : 0.1818
 Neg Pred Value : 1.0000
 Prevalence : 0.0400
 Detection Rate : 0.0400
 Detection Prevalence : 0.2200
 Balanced Accuracy : 0.9062

 'Positive' Class : reptile

2 Classification: Alternative Techniques

2.1 Install packages

```

if(!require(pacman))
  install.packages("pacman")

pacman::p_load(
  C50,           # C5.0 Decision Trees and Rule-Based Models
  caret,         # Classification and Regression Training
  e1071,         # Misc Functions of the Department of Statistics (e1071), TU Wien
  keras,         # R Interface to 'Keras'
  kernlab,       # Kernel-Based Machine Learning Lab
  lattice,       # Trellis Graphics for R

```

```

MASS,          # Support Functions and Datasets for Venables and Ripley's MASS
mlbench,       # Machine Learning Benchmark Problems
nnet,          # Feedforward Neural Networks and Multinomial Log-Linear Models
palmerpenguins, # Palmer Archipelago (Antarctica) Penguin Data
party,         # A Laboratory for Recursive Partytioning
partykit,      # A Toolkit for Recursive Partytioning
randomForest,  # Breiman and Cutler's Random Forests for Classification and Regress
rpart,         # Recursive partitioning models
RWeka,         # R/Weka Interface
scales,        # Scale Functions for Visualization
tidymodels,    # Tidy machine learning framework
tidyverse,     # Tidy data wrangling and visualization
xgboost        # Extreme Gradient Boosting
)

options(digits=3)

```

2.2 Training and Test Data

```

data(Zoo, package="mlbench")
Zoo <- as.data.frame(Zoo)
Zoo |> glimpse()

```

```

Rows: 101
Columns: 17
$ hair      <lgl> TRUE, TRUE, FALSE, TRUE, TRUE, TRUE, TRUE, FALSE, FALSE, TRUE~
$ feathers  <lgl> FALSE, FALSE, FALSE, FALSE, FALSE, FALSE, FALSE, FALSE, FALSE~
$ eggs      <lgl> FALSE, FALSE, TRUE, FALSE, FALSE, FALSE, FALSE, TRUE, TRUE, F~
$ milk      <lgl> TRUE, TRUE, FALSE, TRUE, TRUE, TRUE, TRUE, FALSE, FALSE, TRUE~
$ airborne  <lgl> FALSE, FALSE, FALSE, FALSE, FALSE, FALSE, FALSE, FALSE, FALSE~
$ aquatic   <lgl> FALSE, FALSE, TRUE, FALSE, FALSE, FALSE, FALSE, TRUE, TRUE, F~
$ predator  <lgl> TRUE, FALSE, TRUE, TRUE, TRUE, FALSE, FALSE, FALSE, TRUE, FAL~
$ toothed   <lgl> TRUE, TRUE, TRUE, TRUE, TRUE, TRUE, TRUE, TRUE, TRUE, TRUE, T~
$ backbone  <lgl> TRUE, TRUE, TRUE, TRUE, TRUE, TRUE, TRUE, TRUE, TRUE, TRUE, T~
$ breathes  <lgl> TRUE, TRUE, FALSE, TRUE, TRUE, TRUE, TRUE, FALSE, FALSE, TRUE~
$ venomous  <lgl> FALSE, FALSE, FALSE, FALSE, FALSE, FALSE, FALSE, FALSE, FALSE~
$ fins      <lgl> FALSE, FALSE, TRUE, FALSE, FALSE, FALSE, FALSE, TRUE, TRUE, F~
$ legs      <int> 4, 4, 0, 4, 4, 4, 4, 0, 0, 4, 4, 2, 0, 0, 4, 6, 2, 4, 0, 0, 2~
$ tail      <lgl> FALSE, TRUE, TRUE, FALSE, TRUE, TRUE, TRUE, TRUE, TRUE, FALSE~

```

```
$ domestic <lgl> FALSE, FALSE, FALSE, FALSE, FALSE, FALSE, TRUE, TRUE, FALSE, ~
$ catsize <lgl> TRUE, TRUE, FALSE, TRUE, TRUE, TRUE, TRUE, FALSE, FALSE, FALS~
$ type <fct> mammal, mammal, fish, mammal, mammal, mammal, mammal, fish, f~
```

```
set.seed(123) # for reproducibility
inTrain <- createDataPartition(y = Zoo$type, p = .8)[[1]]
Zoo_train <- dplyr::slice(Zoo, inTrain)
Zoo_test <- dplyr::slice(Zoo, -inTrain)
```

2.3 Fitting Different Classification Models to the Training Data

```
train_index <- createFolds(Zoo_train$type, k = 10)
```

2.3.1 Conditional Inference Tree (Decision Tree)

```
ctreeFit <- Zoo_train |> train(type ~ .,
  method = "ctree",
  data = _,
  tuneLength = 5,
  trControl = trainControl(method = "cv", indexOut = train_index))
ctreeFit
```

Conditional Inference Tree

83 samples

16 predictors

7 classes: 'mammal', 'bird', 'reptile', 'fish', 'amphibian', 'insect', 'mollusc.etal'

No pre-processing

Resampling: Cross-Validated (10 fold)

Summary of sample sizes: 76, 72, 73, 76, 75, 75, ...

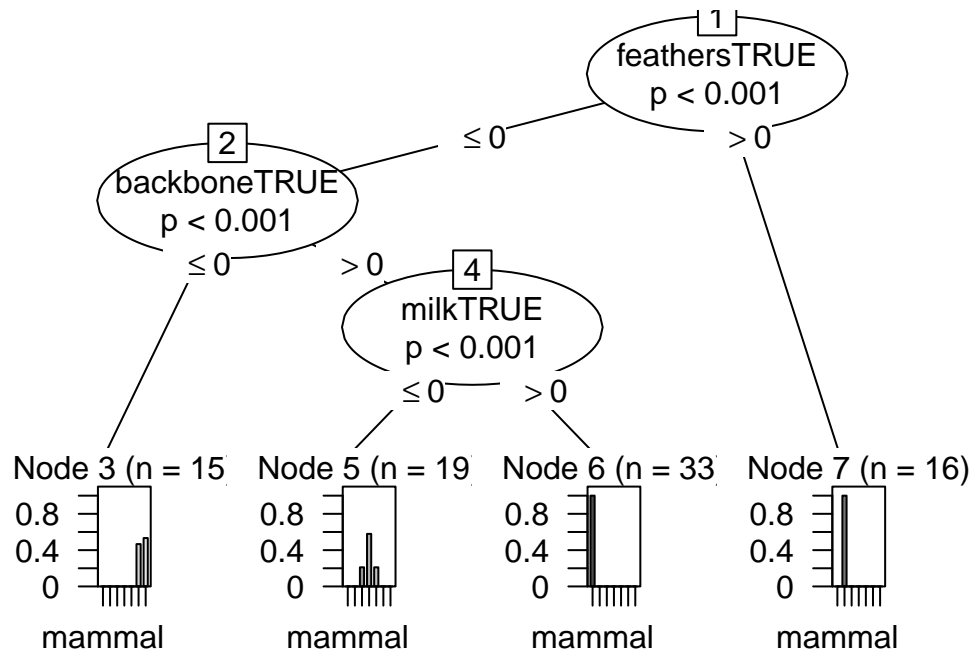
Resampling results across tuning parameters:

mincriterion	Accuracy	Kappa
0.010	0.827	0.772
0.255	0.827	0.772
0.500	0.827	0.772
0.745	0.827	0.772

0.990 0.827 0.772

Accuracy was used to select the optimal model using the largest value.
The final value used for the model was mincriterion = 0.99.

```
plot(ctreeFit$finalModel)
```



2.3.2 C 4.5 Decision Tree

```
C45Fit <- Zoo_train |> train(type ~ .,
  method = "J48",
  data = _,
  tuneLength = 5,
  trControl = trainControl(method = "cv", indexOut = train_index))
C45Fit
```

C4.5-like Trees

83 samples

16 predictors

7 classes: 'mammal', 'bird', 'reptile', 'fish', 'amphibian', 'insect', 'mollusc.et.al'

No pre-processing
 Resampling: Cross-Validated (10 fold)
 Summary of sample sizes: 76, 75, 73, 76, 74, 74, ...
 Resampling results across tuning parameters:

C	M	Accuracy	Kappa
0.010	1	0.975	0.967
0.010	2	0.965	0.954
0.010	3	0.953	0.940
0.010	4	0.959	0.948
0.010	5	0.970	0.962
0.133	1	1.000	1.000
0.133	2	0.976	0.968
0.133	3	0.965	0.954
0.133	4	0.959	0.948
0.133	5	0.970	0.962
0.255	1	1.000	1.000
0.255	2	0.976	0.968
0.255	3	0.965	0.954
0.255	4	0.959	0.948
0.255	5	0.970	0.962
0.378	1	1.000	1.000
0.378	2	0.976	0.968
0.378	3	0.965	0.954
0.378	4	0.959	0.948
0.378	5	0.970	0.962
0.500	1	1.000	1.000
0.500	2	0.976	0.968
0.500	3	0.965	0.954
0.500	4	0.959	0.948
0.500	5	0.970	0.962

Accuracy was used to select the optimal model using the largest value.
 The final values used for the model were C = 0.133 and M = 1.

```
C45Fit$finalModel
```

J48 pruned tree

```
feathersTRUE <= 0
```

```

|   milkTRUE <= 0
|   |   backboneTRUE <= 0
|   |   |   predatorTRUE <= 0
|   |   |   |   legs <= 2: mollusc.et.al (1.0)
|   |   |   |   legs > 2: insect (6.0)
|   |   |   |   predatorTRUE > 0: mollusc.et.al (8.0/1.0)
|   |   |   backboneTRUE > 0
|   |   |   |   finsTRUE <= 0
|   |   |   |   |   aquaticTRUE <= 0: reptile (3.0)
|   |   |   |   |   aquaticTRUE > 0
|   |   |   |   |   |   eggsTRUE <= 0: reptile (1.0)
|   |   |   |   |   |   eggsTRUE > 0: amphibian (4.0)
|   |   |   |   |   |   finsTRUE > 0: fish (11.0)
|   milkTRUE > 0: mammal (33.0)
feathersTRUE > 0: bird (16.0)

```

Number of Leaves : 9

Size of the tree : 17

2.3.3 K-Nearest Neighbors

```

knnFit <- Zoo_train |> train(type ~ .,
  method = "knn",
  data = _,
  preProcess = "scale",
  tuneLength = 5,
  tuneGrid=data.frame(k = 1:10),
  trControl = trainControl(method = "cv", indexOut = train_index))
knnFit

```

k-Nearest Neighbors

83 samples

16 predictors

7 classes: 'mammal', 'bird', 'reptile', 'fish', 'amphibian', 'insect', 'mollusc.et.al'

Pre-processing: scaled (16)

Resampling: Cross-Validated (10 fold)

Summary of sample sizes: 77, 74, 75, 75, 74, 74, ...

Resampling results across tuning parameters:

k	Accuracy	Kappa
1	1.000	1.000
2	0.965	0.954
3	0.963	0.951
4	0.942	0.925
5	0.941	0.921
6	0.963	0.951
7	0.963	0.951
8	0.941	0.921
9	0.908	0.883
10	0.918	0.892

Accuracy was used to select the optimal model using the largest value.
The final value used for the model was $k = 1$.

```
knnFit$finalModel
```

1-nearest neighbor model

Training set outcome distribution:

mammal	bird	reptile	fish	amphibian
33	16	4	11	4
insect	mollusc.et.al			
7	8			

2.3.4 PART (Rule-based classifier)

```
rulesFit <- Zoo_train |> train(type ~ .,
  method = "PART",
  data = _,
  tuneLength = 5,
  trControl = trainControl(method = "cv", indexOut = train_index))
rulesFit
```

Rule-Based Classifier

83 samples

16 predictors

7 classes: 'mammal', 'bird', 'reptile', 'fish', 'amphibian', 'insect', 'mollusc.et.al'

No pre-processing
 Resampling: Cross-Validated (10 fold)
 Summary of sample sizes: 73, 74, 76, 76, 75, 74, ...
 Resampling results across tuning parameters:

threshold	pruned	Accuracy	Kappa
0.010	yes	0.979	0.973
0.010	no	0.979	0.973
0.133	yes	0.990	0.987
0.133	no	0.979	0.973
0.255	yes	0.990	0.987
0.255	no	0.979	0.973
0.378	yes	0.990	0.987
0.378	no	0.979	0.973
0.500	yes	0.990	0.987
0.500	no	0.979	0.973

Accuracy was used to select the optimal model using the largest value.
 The final values used for the model were threshold = 0.5 and pruned = yes.

```
rulesFit$finalModel
```

PART decision list

```
feathersTRUE <= 0 AND
milkTRUE > 0: mammal (33.0)

feathersTRUE > 0: bird (16.0)

backboneTRUE <= 0 AND
airborneTRUE <= 0 AND
predatorTRUE > 0: mollusc.et.al (7.0)

backboneTRUE > 0 AND
finsTRUE > 0: fish (11.0)

backboneTRUE <= 0: insect (8.0/1.0)

aquaticTRUE > 0: amphibian (5.0/1.0)
```

```
: reptile (3.0)
```

```
Number of Rules : 7
```

2.3.5 Linear Support Vector Machines

```
svmFit <- Zoo_train |> train(type ~.,  
  method = "svmLinear",  
  data = _,  
  tuneLength = 5,  
  trControl = trainControl(method = "cv", indexOut = train_index))  
svmFit
```

Support Vector Machines with Linear Kernel

83 samples

16 predictors

7 classes: 'mammal', 'bird', 'reptile', 'fish', 'amphibian', 'insect', 'mollusc.et.al'

No pre-processing

Resampling: Cross-Validated (10 fold)

Summary of sample sizes: 74, 74, 77, 75, 74, 77, ...

Resampling results:

Accuracy	Kappa
1	1

Tuning parameter 'C' was held constant at a value of 1

```
svmFit$finalModel
```

Support Vector Machine object of class "ksvm"

SV type: C-svc (classification)

parameter : cost C = 1

Linear (vanilla) kernel function.

Number of Support Vectors : 39

Objective Function Value : -0.143 -0.217 -0.15 -0.175 -0.0934 -0.0974 -0.292 -0.0835 -0.154

Training error : 0

2.3.6 Random Forest

```
randomForestFit <- Zoo_train |> train(type ~ .,  
  method = "rf",  
  data = _,  
  tuneLength = 5,  
  trControl = trainControl(method = "cv", indexOut = train_index))  
randomForestFit
```

Random Forest

83 samples

16 predictors

7 classes: 'mammal', 'bird', 'reptile', 'fish', 'amphibian', 'insect', 'mollusc.et.al'

No pre-processing

Resampling: Cross-Validated (10 fold)

Summary of sample sizes: 75, 76, 75, 76, 74, 73, ...

Resampling results across tuning parameters:

mtry	Accuracy	Kappa
2	1	1
5	1	1
9	1	1
12	1	1
16	1	1

Accuracy was used to select the optimal model using the largest value.

The final value used for the model was mtry = 2.

```
randomForestFit$finalModel
```

Call:

```
randomForest(x = x, y = y, mtry = param$mtry)  
Type of random forest: classification
```

Number of trees: 500
No. of variables tried at each split: 2

OOB estimate of error rate: 7.23%

Confusion matrix:

	mammal	bird	reptile	fish	amphibian	insect	mollusc.et.al
mammal	33	0	0	0	0	0	0
bird	0	16	0	0	0	0	0
reptile	0	1	0	2	1	0	0
fish	0	0	0	11	0	0	0
amphibian	0	0	0	0	4	0	0
insect	0	0	0	0	0	7	0
mollusc.et.al	1	0	0	0	0	1	6

class.error

mammal	0.00
bird	0.00
reptile	1.00
fish	0.00
amphibian	0.00
insect	0.00
mollusc.et.al	0.25

2.3.7 Gradient Boosted Decision Trees (xgboost)

```
xgboostFit <- Zoo_train |> train(type ~ .,  
  method = "xgbTree",  
  data = _,  
  tuneLength = 5,  
  trControl = trainControl(method = "cv", indexOut = train_index),  
  tuneGrid = expand.grid(  
    nrounds = 20,  
    max_depth = 3,  
    colsample_bytree = .6,  
    eta = 0.1,  
    gamma=0,  
    min_child_weight = 1,  
    subsample = .5  
  ))  
xgboostFit
```

eXtreme Gradient Boosting

83 samples
16 predictors
7 classes: 'mammal', 'bird', 'reptile', 'fish', 'amphibian', 'insect', 'mollusc.et.al'

No pre-processing
Resampling: Cross-Validated (10 fold)
Summary of sample sizes: 77, 73, 74, 75, 75, 75, ...
Resampling results:

Accuracy	Kappa
0.973	0.964

Tuning parameter 'nrounds' was held constant at a value of 20
Tuning
held constant at a value of 1
Tuning parameter 'subsample' was held
constant at a value of 0.5

```
xgboostFit$finalModel
```

```
##### xgb.Booster
raw: 112.4 Kb
call:
  xgboost::xgb.train(params = list(eta = param$eta, max_depth = param$max_depth,
    gamma = param$gamma, colsample_bytree = param$colsample_bytree,
    min_child_weight = param$min_child_weight, subsample = param$subsample),
    data = x, nrounds = param$nrounds, num_class = length(lev),
    objective = "multi:softprob")
params (as set within xgb.train):
  eta = "0.1", max_depth = "3", gamma = "0", colsample_bytree = "0.6", min_child_weight = "1"
xgb.attributes:
  niter
callbacks:
  cb.print.evaluation(period = print_every_n)
# of features: 16
niter: 20
nfeatures : 16
xNames : hairTRUE feathersTRUE eggsTRUE milkTRUE airborneTRUE aquaticTRUE predatorTRUE tooth
problemType : Classification
tuneValue :
```

```

      nrounds max_depth eta gamma colsample_bytree min_child_weight subsample
1         20         3 0.1    0             0.6             1         0.5
obsLevels : mammal bird reptile fish amphibian insect mollusc.et.al
param :
  list()

```

2.3.8 Artificial Neural Network

```

nnetFit <- Zoo_train |> train(type ~ .,
  method = "nnet",
  data = _,
  tuneLength = 5,
  trControl = trainControl(method = "cv", indexOut = train_index),
  trace = FALSE)
nnetFit

```

Neural Network

83 samples

16 predictors

7 classes: 'mammal', 'bird', 'reptile', 'fish', 'amphibian', 'insect', 'mollusc.et.al'

No pre-processing

Resampling: Cross-Validated (10 fold)

Summary of sample sizes: 75, 74, 74, 74, 74, 75, ...

Resampling results across tuning parameters:

size	decay	Accuracy	Kappa
1	0e+00	0.776	0.681
1	1e-04	0.789	0.709
1	1e-03	0.911	0.882
1	1e-02	0.832	0.781
1	1e-01	0.722	0.621
3	0e+00	0.963	0.950
3	1e-04	0.976	0.968
3	1e-03	0.986	0.979
3	1e-02	0.986	0.981
3	1e-01	0.976	0.968
5	0e+00	0.965	0.953
5	1e-04	0.986	0.981
5	1e-03	0.986	0.981

5	1e-02	0.986	0.981
5	1e-01	0.986	0.981
7	0e+00	0.976	0.968
7	1e-04	0.986	0.981
7	1e-03	0.986	0.981
7	1e-02	0.986	0.981
7	1e-01	0.986	0.981
9	0e+00	0.986	0.981
9	1e-04	0.986	0.981
9	1e-03	0.986	0.981
9	1e-02	0.986	0.981
9	1e-01	0.986	0.981

Accuracy was used to select the optimal model using the largest value.
The final values used for the model were size = 3 and decay = 0.01.

```
nnetFit$finalModel
```

a 16-3-7 network with 79 weights

inputs: hairTRUE feathersTRUE eggsTRUE milkTRUE airborneTRUE aquaticTRUE predatorTRUE toothed

output(s): .outcome

options were - softmax modelling decay=0.01

2.4 Comparing Models

```
resamps <- resamples(list(
  ctree = ctreeFit,
  C45 = C45Fit,
  SVM = svmFit,
  KNN = knnFit,
  rules = rulesFit,
  randomForest = randomForestFit,
  xgboost = xgboostFit,
  NeuralNet = nnetFit
))
resamps
```

Call:

```
resamples.default(x = list(ctree = ctreeFit, C45 = C45Fit, SVM = svmFit, KNN
```

```
= knnFit, rules = rulesFit, randomForest = randomForestFit, xgboost
= xgboostFit, NeuralNet = nnetFit))
```

Models: ctree, C45, SVM, KNN, rules, randomForest, xgboost, NeuralNet
 Number of resamples: 10
 Performance metrics: Accuracy, Kappa
 Time estimates for: everything, final model fit

```
summary(resamps)
```

Call:

```
summary.resamples(object = resamps)
```

Models: ctree, C45, SVM, KNN, rules, randomForest, xgboost, NeuralNet
 Number of resamples: 10

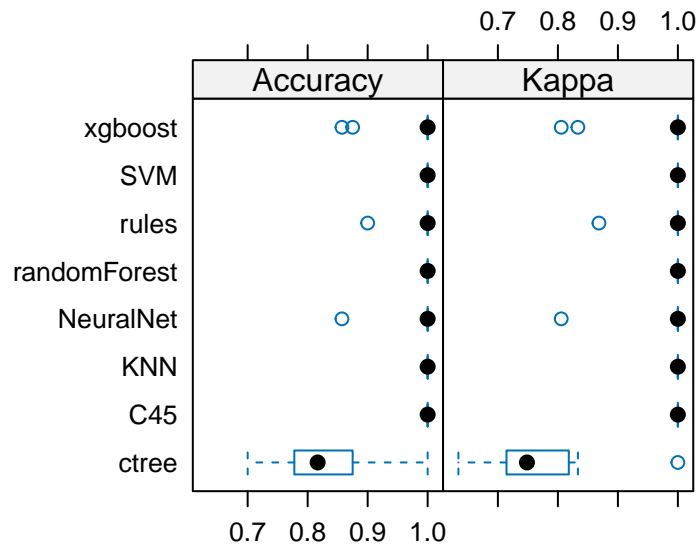
Accuracy

	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.	NA's
ctree	0.700	0.778	0.817	0.827	0.871	1	0
C45	1.000	1.000	1.000	1.000	1.000	1	0
SVM	1.000	1.000	1.000	1.000	1.000	1	0
KNN	1.000	1.000	1.000	1.000	1.000	1	0
rules	0.900	1.000	1.000	0.990	1.000	1	0
randomForest	1.000	1.000	1.000	1.000	1.000	1	0
xgboost	0.857	1.000	1.000	0.973	1.000	1	0
NeuralNet	0.857	1.000	1.000	0.986	1.000	1	0

Kappa

	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.	NA's
ctree	0.634	0.715	0.748	0.772	0.815	1	0
C45	1.000	1.000	1.000	1.000	1.000	1	0
SVM	1.000	1.000	1.000	1.000	1.000	1	0
KNN	1.000	1.000	1.000	1.000	1.000	1	0
rules	0.868	1.000	1.000	0.987	1.000	1	0
randomForest	1.000	1.000	1.000	1.000	1.000	1	0
xgboost	0.806	1.000	1.000	0.964	1.000	1	0
NeuralNet	0.806	1.000	1.000	0.981	1.000	1	0

```
library(lattice)
bwplot(resamps, layout = c(3, 1))
```



```
difs <- diff(resamps)
difs
```

Call:
diff.resamples(x = resamps)

Models: ctree, C45, SVM, KNN, rules, randomForest, xgboost, NeuralNet
Metrics: Accuracy, Kappa
Number of differences: 28
p-value adjustment: bonferroni

```
summary(difs)
```

Call:
summary.diff.resamples(object = difs)

p-value adjustment: bonferroni
Upper diagonal: estimates of the difference
Lower diagonal: p-value for H0: difference = 0

Accuracy		ctree	C45	SVM	KNN	rules	randomForest	xgboost
ctree			-0.17262	-0.17262	-0.17262	-0.16262	-0.17262	-0.14583
C45		0.00193		0.00000	0.00000	0.01000	0.00000	0.02679

SVM	0.00193	NA		0.00000	0.01000	0.00000	0.02679
KNN	0.00193	NA	NA		0.01000	0.00000	0.02679
rules	0.00376	1.00000	1.00000	1.00000		-0.01000	0.01679
randomForest	0.00193	NA	NA	NA	1.00000		0.02679
xgboost	0.05129	1.00000	1.00000	1.00000	1.00000	1.00000	
NeuralNet	0.01405	1.00000	1.00000	1.00000	1.00000	1.00000	1.00000

NeuralNet

ctree	-0.15833
C45	0.01429
SVM	0.01429
KNN	0.01429
rules	0.00429
randomForest	0.01429
xgboost	-0.01250
NeuralNet	

Kappa

	ctree	C45	SVM	KNN	rules	randomForest	xgboost
ctree		-0.22840	-0.22840	-0.22840	-0.21524	-0.22840	-0.19229
C45	0.00116		0.00000	0.00000	0.01316	0.00000	0.03611
SVM	0.00116	NA		0.00000	0.01316	0.00000	0.03611
KNN	0.00116	NA	NA		0.01316	0.00000	0.03611
rules	0.00238	1.00000	1.00000	1.00000		-0.01316	0.02295
randomForest	0.00116	NA	NA	NA	1.00000		0.03611
xgboost	0.04216	1.00000	1.00000	1.00000	1.00000	1.00000	
NeuralNet	0.01055	1.00000	1.00000	1.00000	1.00000	1.00000	1.00000

NeuralNet

ctree	-0.20895
C45	0.01944
SVM	0.01944
KNN	0.01944
rules	0.00629
randomForest	0.01944
xgboost	-0.01667
NeuralNet	

2.5 Applying the Chosen Model to the Test Data

```
pr <- predict(randomForestFit, Zoo_test)
pr
```

```
[1] mammal      mammal      mammal      fish        fish
[6] bird        bird        mammal      mammal      mammal
[11] mammal      mollusc.et.al reptile     mammal      bird
[16] mollusc.et.al bird        insect
Levels: mammal bird reptile fish amphibian insect mollusc.et.al
```

```
confusionMatrix(pr, reference = Zoo_test$type)
```

Confusion Matrix and Statistics

	Reference						
Prediction	mammal	bird	reptile	fish	amphibian	insect	mollusc.et.al
mammal	8	0	0	0	0	0	0
bird	0	4	0	0	0	0	0
reptile	0	0	1	0	0	0	0
fish	0	0	0	2	0	0	0
amphibian	0	0	0	0	0	0	0
insect	0	0	0	0	0	1	0
mollusc.et.al	0	0	0	0	0	0	2

Overall Statistics

```
Accuracy : 1
 95% CI : (0.815, 1)
No Information Rate : 0.444
P-Value [Acc > NIR] : 4.58e-07
```

```
Kappa : 1
```

```
McNemar's Test P-Value : NA
```

Statistics by Class:

	Class: mammal	Class: bird	Class: reptile	Class: fish
Sensitivity	1.000	1.000	1.0000	1.000

Specificity	1.000	1.000	1.0000	1.000
Pos Pred Value	1.000	1.000	1.0000	1.000
Neg Pred Value	1.000	1.000	1.0000	1.000
Prevalence	0.444	0.222	0.0556	0.111
Detection Rate	0.444	0.222	0.0556	0.111
Detection Prevalence	0.444	0.222	0.0556	0.111
Balanced Accuracy	1.000	1.000	1.0000	1.000
Class: amphibian Class: insect Class: mollusc.et.al				
Sensitivity	NA	1.0000		1.000
Specificity	1	1.0000		1.000
Pos Pred Value	NA	1.0000		1.000
Neg Pred Value	NA	1.0000		1.000
Prevalence	0	0.0556		0.111
Detection Rate	0	0.0556		0.111
Detection Prevalence	0	0.0556		0.111
Balanced Accuracy	NA	1.0000		1.000

2.6 Comparing Decision Boundaries of Popular Classification Techniques

```
library(scales)
library(tidyverse)
library(ggplot2)
library(caret)

decisionplot <- function(model, data, class_var,
  predict_type = c("class", "prob"), resolution = 3 * 72) {
  # resolution is set to 72 dpi if the image is rendered 3 inches wide.

  y <- data |> pull(class_var)
  x <- data |> dplyr::select(-all_of(class_var))

  # resubstitution accuracy
  prediction <- predict(model, x, type = predict_type[1])
  # LDA returns a list
  if(is.list(prediction)) prediction <- prediction$class
  prediction <- factor(prediction, levels = levels(y))

  cm <- confusionMatrix(data = prediction,
    reference = y)
  acc <- cm$overall["Accuracy"]
}
```



```

# evaluate model on a grid
r <- sapply(x[, 1:2], range, na.rm = TRUE)
xs <- seq(r[1,1], r[2,1], length.out = resolution)
ys <- seq(r[1,2], r[2,2], length.out = resolution)
g <- cbind(rep(xs, each = resolution), rep(ys, time = resolution))
colnames(g) <- colnames(r)
g <- as_tibble(g)

### guess how to get class labels from predict
### (unfortunately not very consistent between models)
cl <- predict(model, g, type = predict_type[1])

# LDA returns a list
prob <- NULL
if(is.list(cl)) {
  prob <- cl$posterior
  cl <- cl$class
} else
  if(!is.na(predict_type[2]))
    try(prob <- predict(model, g, type = predict_type[2]))

# we visualize the difference in probability/score between the
# winning class and the second best class.
# don't use probability if predict for the classifier does not support it.
max_prob <- 1
if(!is.null(prob))
  try({
    max_prob <- t(apply(prob, MARGIN = 1, sort, decreasing = TRUE))
    max_prob <- max_prob[,1] - max_prob[,2]
  }, silent = TRUE)

cl <- factor(cl, levels = levels(y))

g <- g |> add_column(prediction = cl, probability = max_prob)

ggplot(g, mapping = aes(
  x = .data[[colnames(g)[1]], y = .data[[colnames(g)[2]]]) +
  geom_raster(mapping = aes(fill = prediction, alpha = probability)) +
  geom_contour(mapping = aes(z = as.numeric(prediction)),
    bins = length(levels(cl)), linewidth = .5, color = "black") +
  geom_point(data = data, mapping = aes(

```

```

    x = .data[[colnames(data)[1]]],
    y = .data[[colnames(data)[2]]],
    shape = .data[[class_var]], alpha = .7) +
scale_alpha_continuous(range = c(0,1), limits = c(0,1), guide = "none") +
labs(subtitle = paste("Training accuracy:", round(acc, 2))) +
theme_minimal(base_size = 14)
}

```

2.6.1 Penguins Dataset

```

set.seed(1000)
data("penguins")
penguins <- as_tibble(penguins) |>
  drop_na()

### Three classes
### (note: MASS also has a select function which hides dplyr's select)
x <- penguins |> dplyr::select(bill_length_mm, bill_depth_mm, species)
x

```

```

# A tibble: 333 x 3
  bill_length_mm bill_depth_mm species
      <dbl>         <dbl> <fct>
1         39.1         18.7 Adelie
2         39.5         17.4 Adelie
3         40.3          18   Adelie
4         36.7         19.3 Adelie
5         39.3         20.6 Adelie
6         38.9         17.8 Adelie
7         39.2         19.6 Adelie
8         41.1         17.6 Adelie
9         38.6         21.2 Adelie
10        34.6         21.1 Adelie
# i 323 more rows


```

```

ggplot(x, aes(x = bill_length_mm, y = bill_depth_mm, fill = species)) +
  stat_density_2d(geom = "polygon", aes(alpha = after_stat(level))) +
  geom_point() +
  theme_minimal(base_size = 14) +

```

```
labs(x = "Bill length (mm)",
     y = "Bill depth (mm)",
     fill = "Species",
     alpha = "Density")
```




2.6.1.1 K-Nearest Neighbors Classifier

```
model <- x |> caret::knn3(species ~ ., data = _, k = 1)
decisionplot(model, x, class_var = "species") +
  labs(title = "kNN (1 neighbor)",
       x = "Bill length (mm)",
       y = "Bill depth (mm)",
       shape = "Species",
       fill = "Prediction")
```

kNN (1 neighbor)

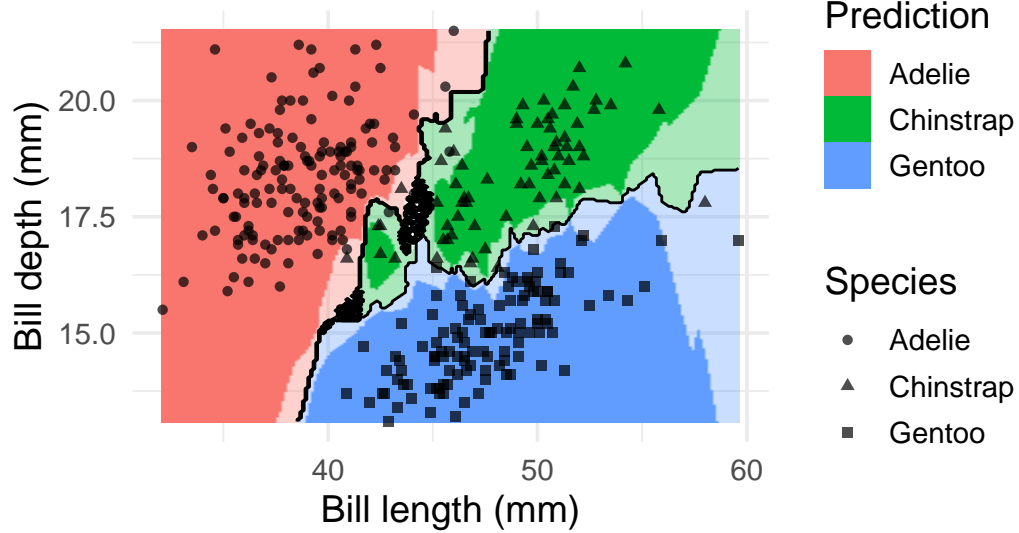
Training accuracy: 1



```
model <- x |> caret::knn3(species ~ ., data = _, k = 3)
decisionplot(model, x, class_var = "species") +
  labs(title = "kNN (3 neighbor)",
        x = "Bill length (mm)",
        y = "Bill depth (mm)",
        shape = "Species",
        fill = "Prediction")
```

kNN (3 neighbor)


Training accuracy: 0.98



```
model <- x |> caret::knn3(species ~ ., data = _, k = 9)
decisionplot(model, x, class_var = "species") +
  labs(title = "kNN (9 neighbor)",
        x = "Bill length (mm)",
        y = "Bill depth (mm)",
        shape = "Species",
        fill = "Prediction")
```

kNN (9 neighbor)

Training accuracy: 0.96




2.6.1.2 Naive Bayes Classifier

```
model <- x |> e1071::naiveBayes(species ~ ., data = _)
decisionplot(model, x, class_var = "species",
             predict_type = c("class", "raw")) +
labs(title = "Naive Bayes",
     x = "Bill length (mm)",
     y = "Bill depth (mm)",
     shape = "Species",
     fill = "Prediction")
```

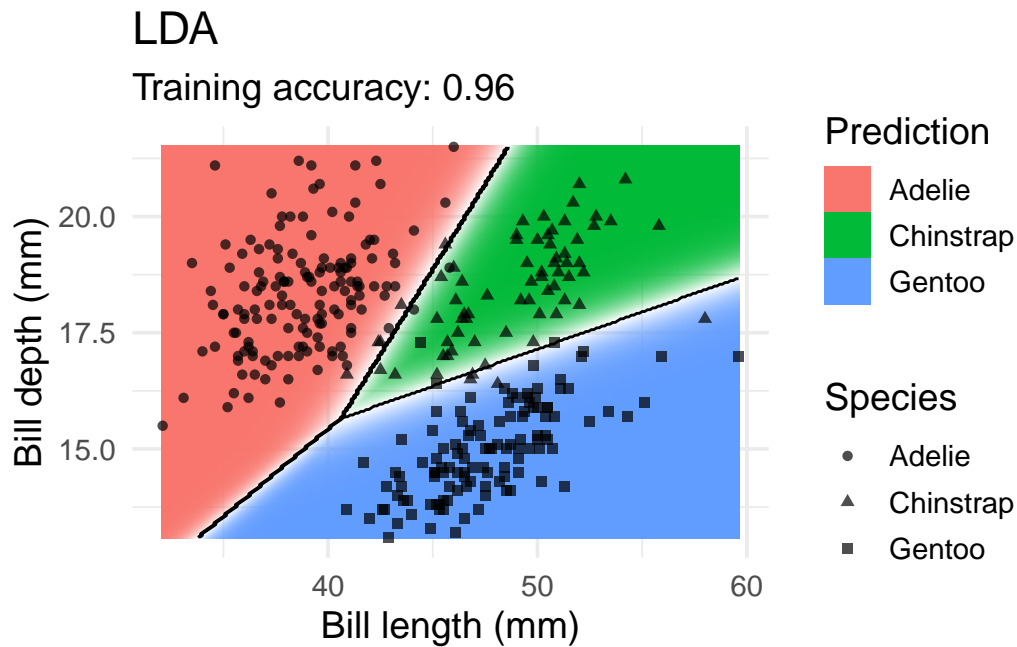
Naive Bayes

Training accuracy: 0.93



2.6.1.3 Linear Discriminant Analysis

```
model <- x |> MASS::lda(species ~ ., data = _)
decisionplot(model, x, class_var = "species") +
  labs(title = "LDA",
        x = "Bill length (mm)",
        y = "Bill depth (mm)",
        shape = "Species",
        fill = "Prediction")
```



2.6.1.4 Multinomial Logistic Regression (implemented in nnet)


```
model <- x |> nnet::multinom(species ~., data = _)
```

```
# weights: 12 (6 variable)
initial value 365.837892
iter 10 value 26.650783
iter 20 value 23.943597
iter 30 value 23.916873
iter 40 value 23.901339
iter 50 value 23.895442
iter 60 value 23.894251
final value 23.892065
converged
```

```
decisionplot(model, x, class_var = "species") +
  labs(title = "Multinomial Logistic Regression",
       x = "Bill length (mm)",
       y = "Bill depth (mm)",
       shape = "Species",
       fill = "Prediction")
```


Multinomial Logistic Regression

Training accuracy: 0.96




2.6.1.5 Decision Trees

```
model <- x |> rpart::rpart(species ~ ., data = _)
decisionplot(model, x, class_var = "species") +
  labs(title = "CART",
        x = "Bill length (mm)",
        y = "Bill depth (mm)",
        shape = "Species",
        fill = "Prediction")
```

CART


Training accuracy: 0.94



```
model <- x |> rpart::rpart(species ~ ., data = _,
  control = rpart.control(cp = 0.001, minsplit = 1))
decisionplot(model, x, class_var = "species") +
  labs(title = "CART (overfitting)",
    x = "Bill length (mm)",
    y = "Bill depth (mm)",
    shape = "Species",
    fill = "Prediction")
```

CART (overfitting)


Training accuracy: 1



```
model <- x |> C50::C5.0(species ~ ., data = _)
decisionplot(model, x, class_var = "species") +
  labs(title = "C5.0",
        x = "Bill length (mm)",
        y = "Bill depth (mm)",
        shape = "Species",
        fill = "Prediction")
```

C5.0

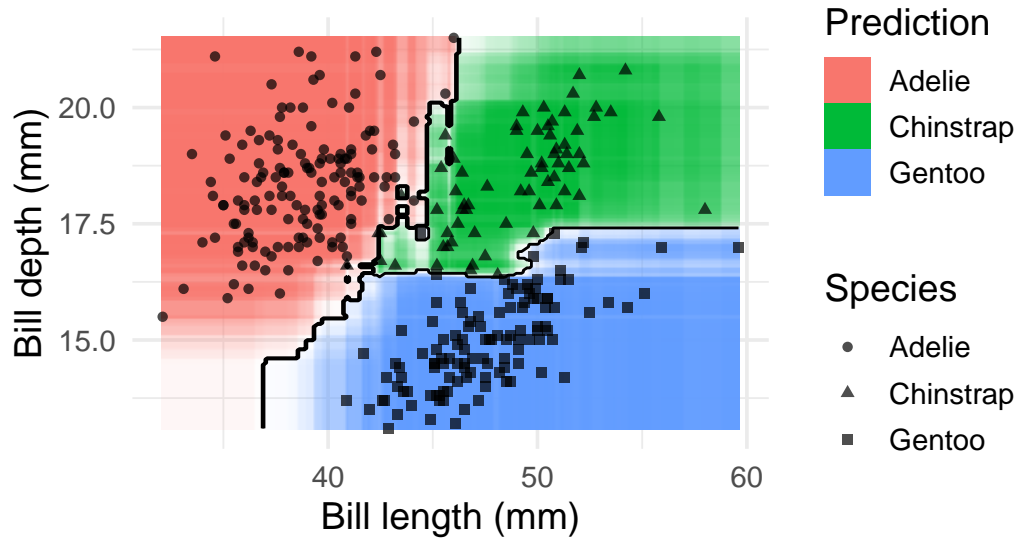
Training accuracy: 0.98



```
model <- x |> randomForest::randomForest(species ~ ., data = _)
decisionplot(model, x, class_var = "species") +
  labs(title = "Random Forest",
        x = "Bill length (mm)",
        y = "Bill depth (mm)",
        shape = "Species",
        fill = "Prediction")
```

Random Forest

Training accuracy: 1




2.6.1.6 SVM

```
model <- x |> e1071::svm(species ~ ., data = _, kernel = "linear")
decisionplot(model, x, class_var = "species") +
  labs(title = "SVM (linear kernel)",
        x = "Bill length (mm)",
        y = "Bill depth (mm)",
        shape = "Species",
        fill = "Prediction")
```

SVM (linear kernel)

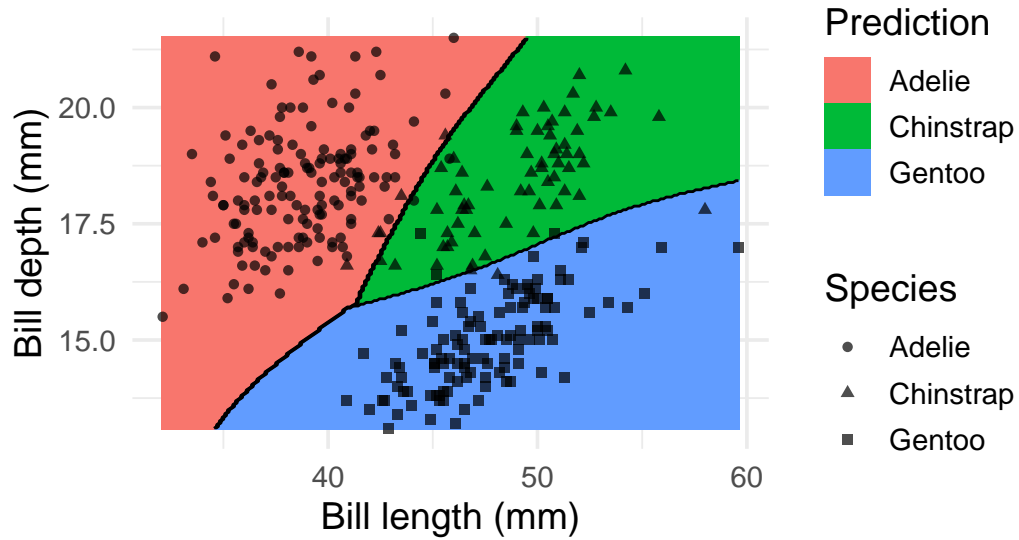
Training accuracy: 0.96



```
model <- x |> e1071::svm(species ~ ., data = _, kernel = "radial")
decisionplot(model, x, class_var = "species") +
  labs(title = "SVM (radial kernel)",
        x = "Bill length (mm)",
        y = "Bill depth (mm)",
        shape = "Species",
        fill = "Prediction")
```

SVM (radial kernel)


Training accuracy: 0.96



```
model <- x |> e1071::svm(species ~ ., data = _, kernel = "polynomial")
decisionplot(model, x, class_var = "species") +
  labs(title = "SVM (polynomial kernel)",
        x = "Bill length (mm)",
        y = "Bill depth (mm)",
        shape = "Species",
        fill = "Prediction")
```

SVM (polynomial kernel)


Training accuracy: 0.91



```
model <- x |> e1071::svm(species ~ ., data = _, kernel = "sigmoid")
decisionplot(model, x, class_var = "species") +
  labs(title = "SVM (sigmoid kernel)",
       x = "Bill length (mm)",
       y = "Bill depth (mm)",
       shape = "Species",
       fill = "Prediction")
```


SVM (sigmoid kernel)

Training accuracy: 0.94




2.6.1.7 Single Layer Feed-forward Neural Networks

```
model <-x |> nnet::nnet(species ~ ., data = _, size = 1, trace = FALSE)
decisionplot(model, x, class_var = "species",
  predict_type = c("class", "raw")) +
  labs(title = "NN (1 neuron)",
    x = "Bill length (mm)",
    y = "Bill depth (mm)",
    shape = "Species",
    fill = "Prediction")
```

Warning: Computation failed in `stat_contour()`
Caused by error in `if (zero_range(range)) ...`:
! missing value where TRUE/FALSE needed

NN (1 neuron)

Training accuracy: 0.44




```
model <-x |> nnet::nnet(species ~ ., data = _, size = 2, trace = FALSE)
decisionplot(model, x, class_var = "species",
  predict_type = c("class", "raw")) +
  labs(title = "NN (2 neurons)",
    x = "Bill length (mm)",
    y = "Bill depth (mm)",
    shape = "Species",
    fill = "Prediction")
```

Warning: Computation failed in `stat_contour()`
Caused by error in `if (zero_range(range)) ...`:
! missing value where TRUE/FALSE needed

NN (2 neurons)


Training accuracy: 0.44



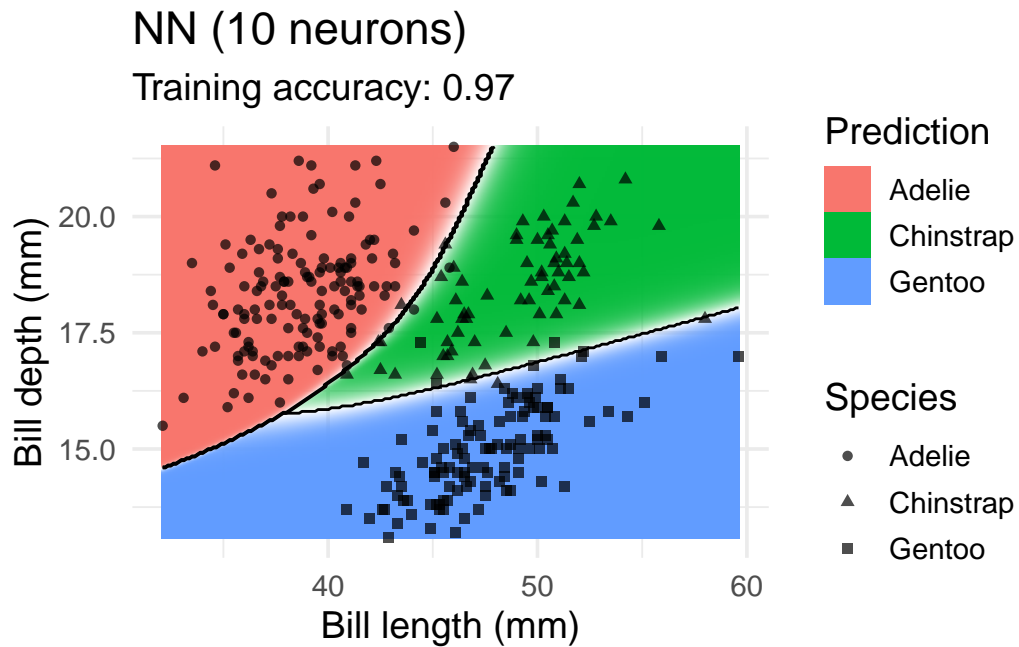
```
model <-x |> nnet::nnet(species ~ ., data = _, size = 4, trace = FALSE)
decisionplot(model, x, class_var = "species",
  predict_type = c("class", "raw")) +
  labs(title = "NN (4 neurons)",
    x = "Bill length (mm)",
    y = "Bill depth (mm)",
    shape = "Species",
    fill = "Prediction")
```

NN (4 neurons)

Training accuracy: 0.95



```
model <-x |> nnet::nnet(species ~ ., data = _, size = 10, trace = FALSE)
decisionplot(model, x, class_var = "species",
  predict_type = c("class", "raw")) +
labs(title = "NN (10 neurons)",
  x = "Bill length (mm)",
  y = "Bill depth (mm)",
  shape = "Species",
  fill = "Prediction")
```



2.6.2 Circle Dataset

```
set.seed(1000)

x <- mlbench::mlbench.circle(500)
###x <- mlbench::mlbench.cassini(500)
###x <- mlbench::mlbench.spirals(500, sd = .1)
###x <- mlbench::mlbench.smiley(500)
x <- cbind(as.data.frame(x$x), factor(x$class))
colnames(x) <- c("x", "y", "class")
x <- as_tibble(x)
x
```

```
# A tibble: 500 x 3
      x      y class
  <dbl> <dbl> <fct>
1 -0.344  0.448  1
2  0.518  0.915  2
3 -0.772 -0.0913 1
4  0.382  0.412  1
5  0.0328 0.438  1
6 -0.865 -0.354  2
```

```


7  0.477  0.640  2
8  0.167 -0.809  2
9 -0.568 -0.281  1
10 -0.488  0.638  2
# i 490 more rows

```

```

ggplot(x, aes(x = x, y = y, color = class)) +
  geom_point() +
  theme_minimal(base_size = 14)

```



2.6.2.1 K-Nearest Neighbors Classifier


```

model <- x |> caret::knn3(class ~ ., data = _, k = 1)
decisionplot(model, x, class_var = "class") +
  labs(title = "kNN (1 neighbor)",
       shape = "Class",
       fill = "Prediction")

```

kNN (1 neighbor)


Training accuracy: 1



```
model <- x |> caret::knn3(class ~ ., data = _, k = 10)
decisionplot(model, x, class_var = "class") +
  labs(title = "kNN (10 neighbor)",
       shape = "Class",
       fill = "Prediction")
```


kNN (10 neighbor)

Training accuracy: 0.98




2.6.2.2 Naive Bayes Classifier

```
model <- x |> e1071::naiveBayes(class ~ ., data = _)
decisionplot(model, x, class_var = "class",
  predict_type = c("class", "raw")) +
  labs(title = "naive Bayes",
    shape = "Class",
    fill = "Prediction")
```



2.6.2.3 Linear Discriminant Analysis

```
model <- x |> MASS::lda(class ~ ., data = _)
decisionplot(model, x, class_var = "class") +
  labs(title = "LDA",
    shape = "Class",
    fill = "Prediction")
```

2.6.2.4 Logistic Regression (implemented in nnet)


```
model <- x |> nnet::multinom(class ~., data = _)
```

```
# weights:  4 (3 variable)
initial  value 346.573590
final    value 346.308371
converged
```

```
decisionplot(model, x, class_var = "class") +
  labs(title = "Multinomial Logistic Regression",
       shape = "Class",
       fill = "Prediction")
```

Multinomial Logistic Regression

Training accuracy: 0.39




2.6.2.5 Decision Trees

```
model <- x |> rpart::rpart(class ~ ., data = _)
decisionplot(model, x, class_var = "class") +
  labs(title = "CART",
       shape = "Class",
       fill = "Prediction")
```

CART


Training accuracy: 0.97



```
model <- x |> rpart::rpart(class ~ ., data = _,  
  control = rpart.control(cp = 0.001, minsplit = 1))  
decisionplot(model, x, class_var = "class") +  
  labs(title = "CART (overfitting)",  
    shape = "Class",  
    fill = "Prediction")
```

CART (overfitting)


Training accuracy: 1




```
model <- x |> C50::C5.0(class ~ ., data = _)
decisionplot(model, x, class_var = "class") +
  labs(title = "C5.0",
        shape = "Class",
        fill = "Prediction")
```

C5.0

Training accuracy: 0.99



```
library(randomForest)
model <- x |> randomForest(class ~ ., data = _)
decisionplot(model, x, class_var = "class") +
  labs(title = "Random Forest",
        shape = "Class",
        fill = "Prediction")
```




2.6.2.6 SVM

```
model <- x |> e1071::svm(class ~ ., data = _, kernel = "linear")
decisionplot(model, x, class_var = "class") +
  labs(title = "SVM (linear kernel)",
        shape = "Class",
        fill = "Prediction")
```

Warning: Computation failed in `stat_contour()`
Caused by error in `if (zero_range(range)) ...`:
! missing value where TRUE/FALSE needed

SVM (linear kernel)


Training accuracy: 0.51



```
model <- x |> e1071::svm(class ~ ., data = _, kernel = "radial")
decisionplot(model, x, class_var = "class") +
  labs(title = "SVM (radial kernel)",
       shape = "Class",
       fill = "Prediction")
```

SVM (radial kernel)


Training accuracy: 0.97



```

model <- x |> e1071::svm(class ~ ., data = _, kernel = "polynomial")
decisionplot(model, x, class_var = "class") +
  labs(title = "SVM (polynomial kernel)",
        shape = "Class",
        fill = "Prediction")

```




```

model <- x |> e1071::svm(class ~ ., data = _, kernel = "sigmoid")
decisionplot(model, x, class_var = "class") +
  labs(title = "SVM (sigmoid kernel)",
        shape = "Class",
        fill = "Prediction")

```

SVM (sigmoid kernel)

Training accuracy: 0.58




2.6.2.7 Single Layer Feed-forward Neural Networks

```
model <-x |> nnet::nnet(class ~ ., data = _, size = 1, trace = FALSE)
decisionplot(model, x, class_var = "class",
  predict_type = c("class")) +
  labs(title = "NN (1 neuron)",
    shape = "Class",
    fill = "Prediction")
```


NN (1 neuron)


Training accuracy: 0.62



```
model <-x |> nnet::nnet(class ~ ., data = _, size = 2, trace = FALSE)
decisionplot(model, x, class_var = "class",
  predict_type = c("class")) +
  labs(title = "NN (2 neurons)",
    shape = "Class",
    fill = "Prediction")
```

NN (2 neurons)


Training accuracy: 0.75



```
model <-x |> nnet::nnet(class ~ ., data = _, size = 4, trace = FALSE)
decisionplot(model, x, class_var = "class",
  predict_type = c("class")) +
  labs(title = "NN (4 neurons)",
    shape = "Class",
    fill = "Prediction")
```

NN (4 neurons)


Training accuracy: 1



```
model <-x |> nnet::nnet(class ~ ., data = _, size = 10, trace = FALSE)
decisionplot(model, x, class_var = "class",
  predict_type = c("class")) +
  labs(title = "NN (10 neurons)",
    shape = "Class",
    fill = "Prediction")
```

NN (10 neurons)

Training accuracy: 1



2.7 More Information on Classification with R

- Package caret: <http://topepo.github.io/caret/index.html>
- Tidymodels (machine learning with tidyverse): <https://www.tidymodels.org/>
- R taskview on machine learning: <http://cran.r-project.org/web/views/MachineLearning.html>